

# CS 246: Artificial Intelligence



Instructor: Br. Tamal Mj

Image credit  
<https://futureoflife.org/>

[slides adapted from Dan Klein, Pieter Abbeel, Sergey Levine & Stuart Russel (University of California, Berkeley)]





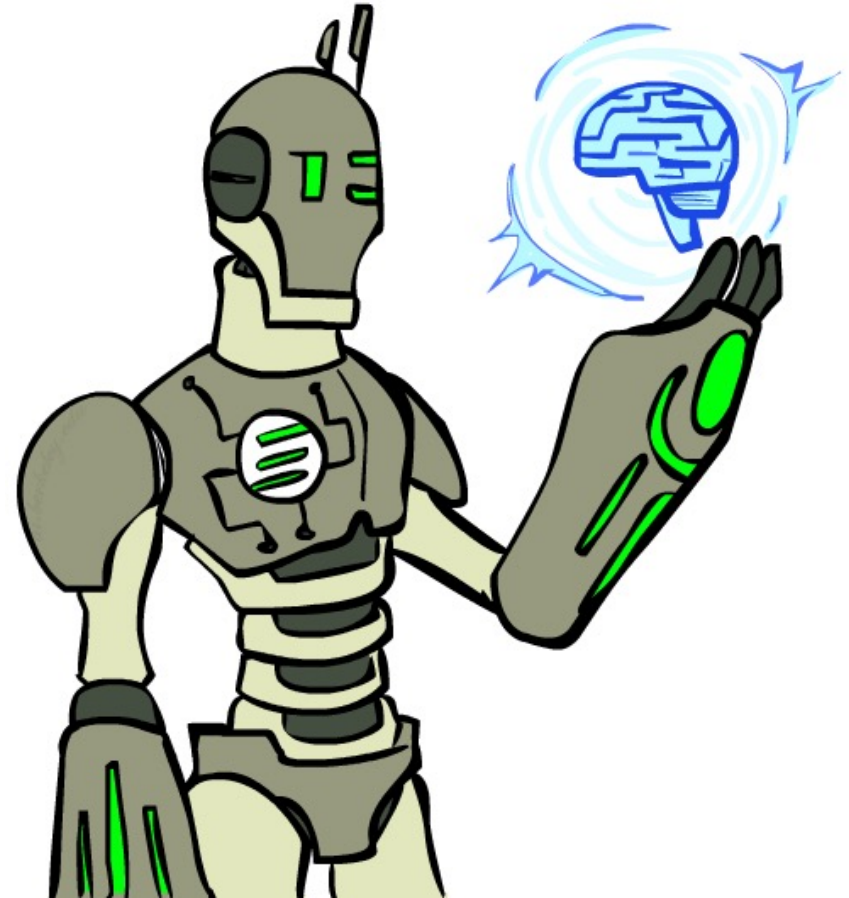
Om Saha Naav[au]-Avatu  
Saha Nau Bhunaktu  
Saha Viiryam Karavaavahai  
Tejasvi Naav[au]-Adhiitam-  
Astu Maa Vidvissaavahai  
Om Shaantih Shaantih  
Shaantih

Om, May we all be protected  
May we all be nourished  
May we work together with great energy  
May our intellect be sharpened (may our study be effective)  
Let there be no Animosity amongst us  
Om, peace (in me), peace (in nature), peace (in divine forces)

# Today

---

- What is artificial intelligence?
- What can AI do?

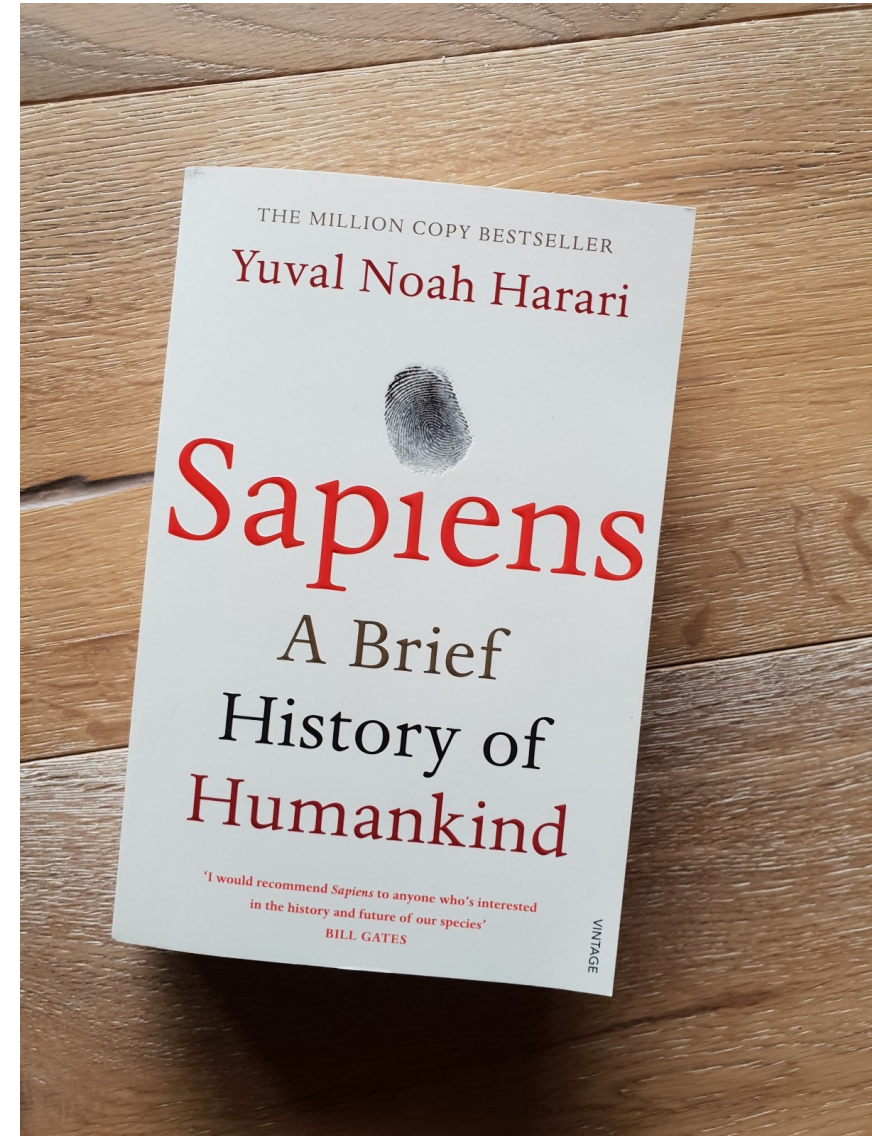




# Homo Sapiens

- Homo Sapiens Man the wise – man the wise – because of our intelligence
- Understanding how we think?
  - Perceive
  - Understand
  - Predict
  - Manipulate

AI goes beyond: Not only understanding but also building intelligent entities



# What is AI?

---

The science of making machines that:



# Rational Decisions

---

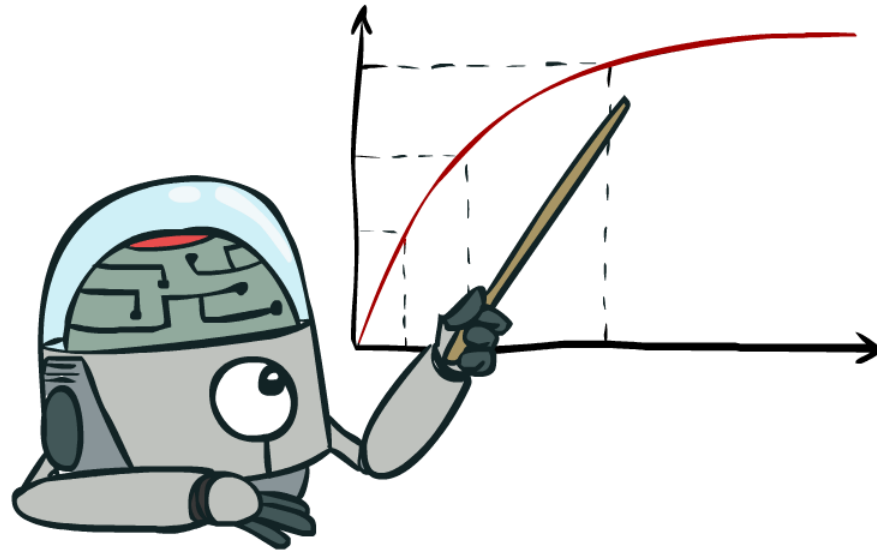
We'll use the term **rational** in a very specific, technical way:

- Rational: maximally achieving pre-defined goals
- Rationality only concerns what decisions are made  
(not the thought process behind them)
- Goals are expressed in terms of the **utility** of outcomes
- Being rational means **maximizing your expected utility**

A better title for this course would be:

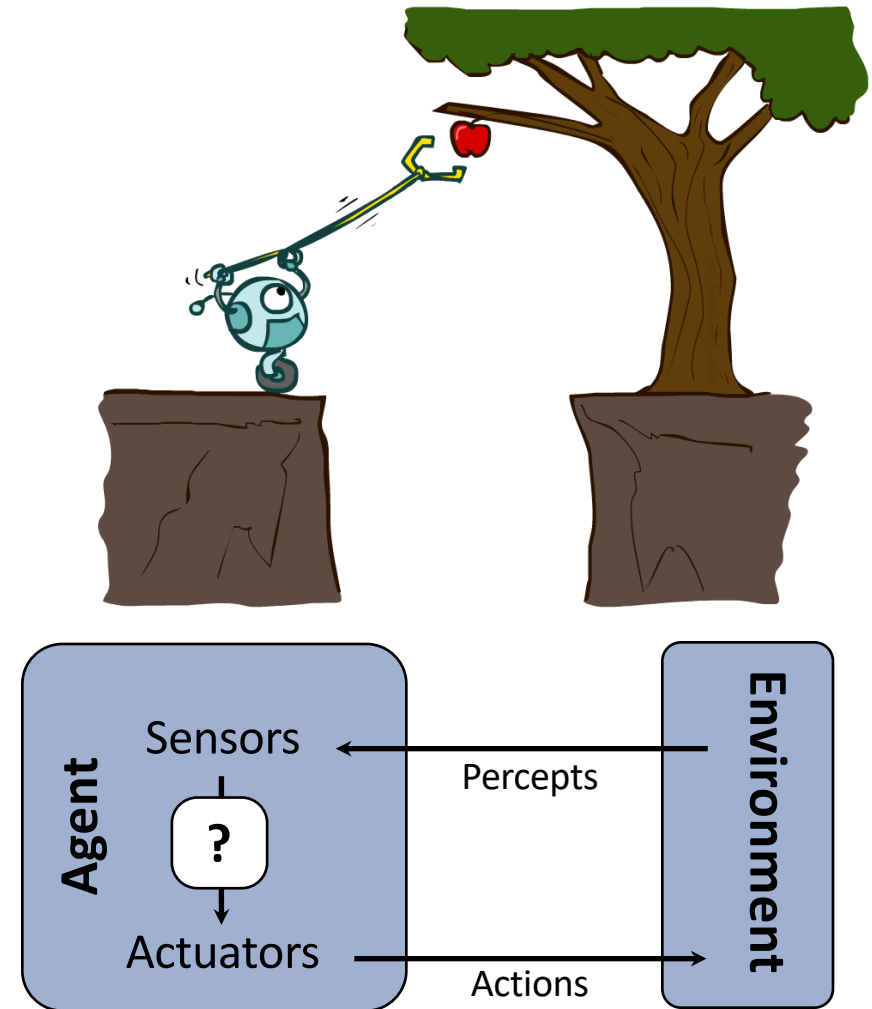
**Computational Rationality**

# Maximize Your Expected Utility



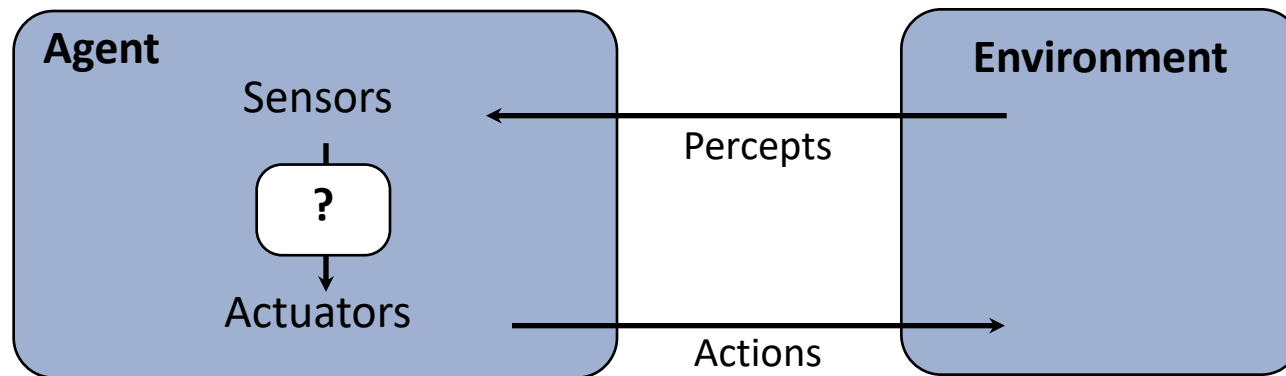
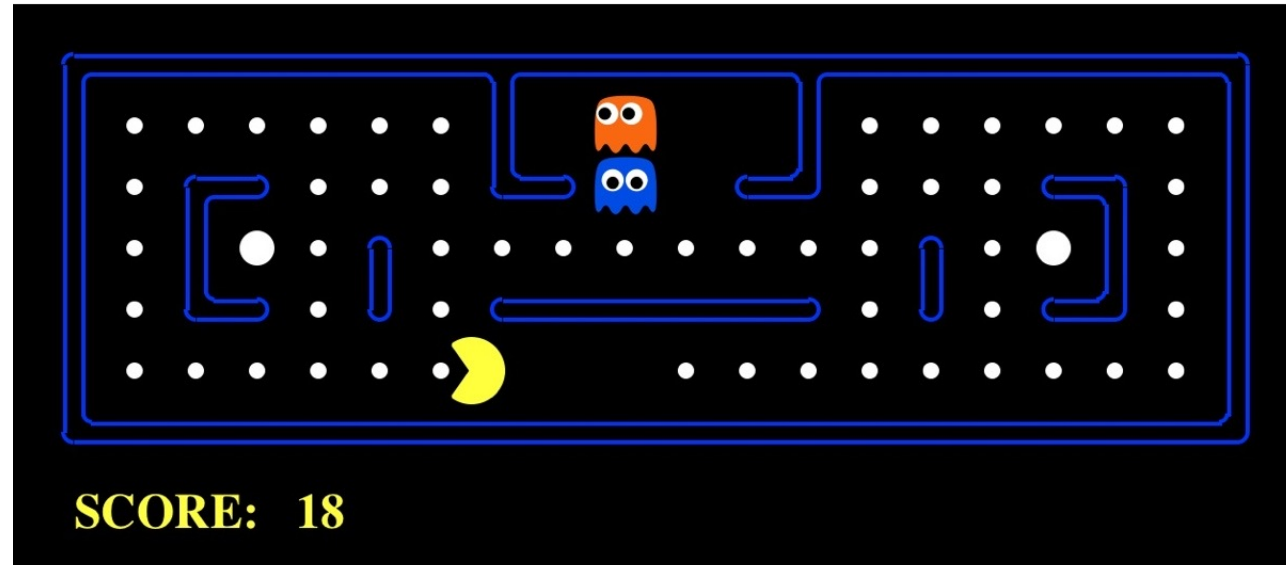
# Designing Rational Agents

- An **agent** is an entity that *perceives* and *acts*.
- A **rational agent** selects actions that maximize its (expected) **utility**.
- Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions
- **This course** is about:
  - General AI techniques for a variety of problem types
  - Learning to recognize when and how a new problem can be solved with an existing technique



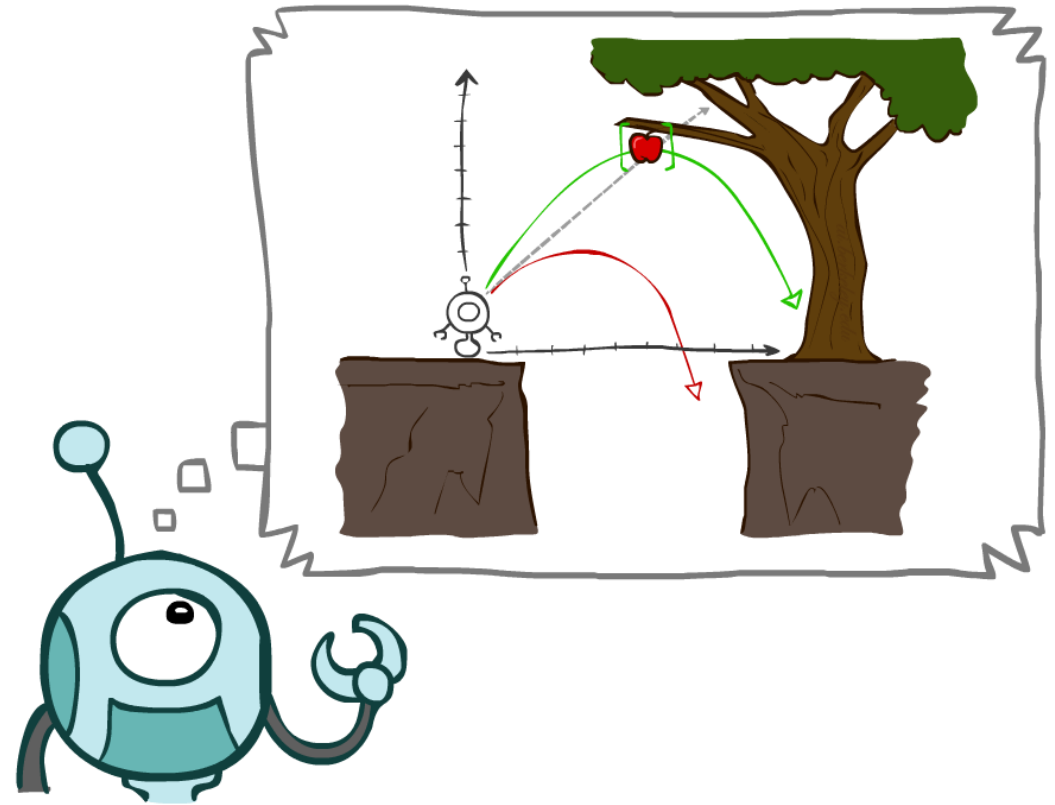


# Pac-Man as an Agent

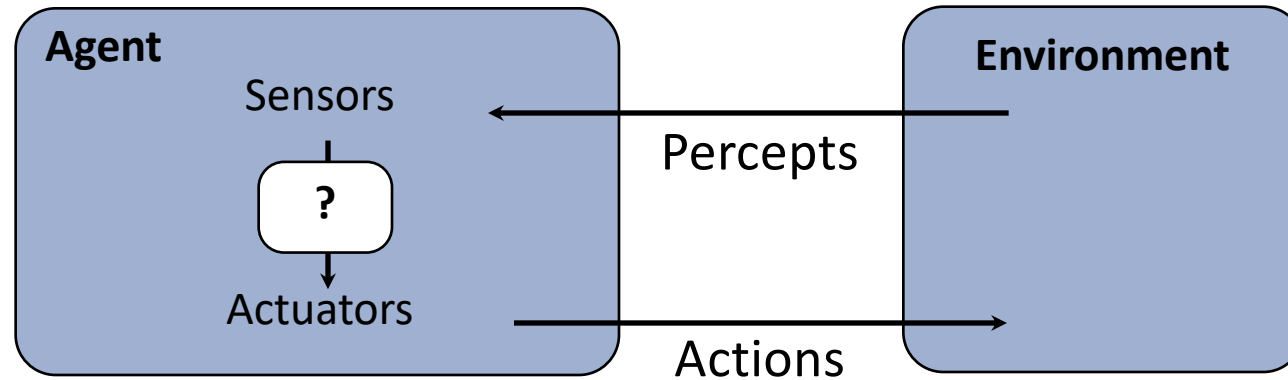


# Next

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
  - Depth-First Search
  - Breadth-First Search
  - Uniform-Cost Search



# Agents and environments



- An agent *perceives* its environment through *sensors* and *acts* upon it through *actuators*



# Rationality

---

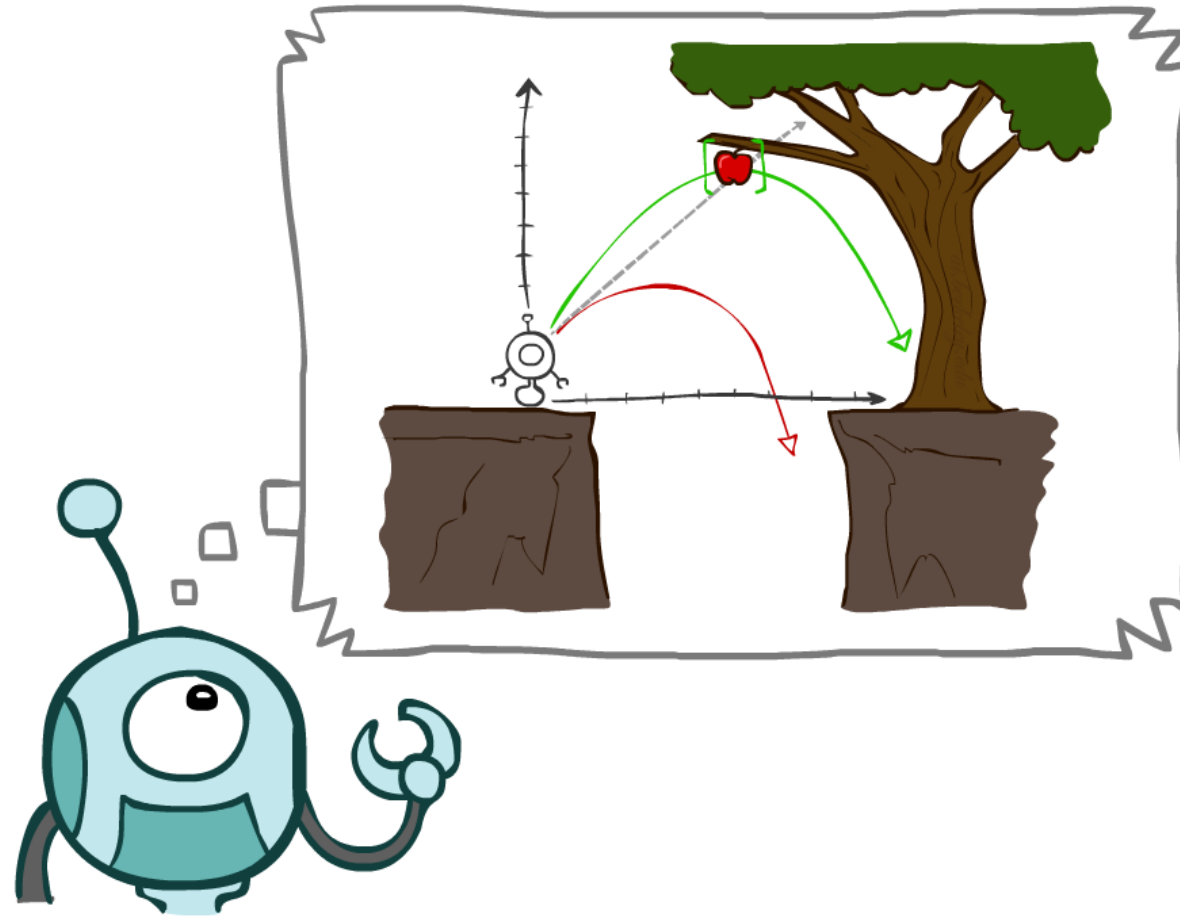
- A *rational agent* chooses actions maximize the *expected* utility
  - Today: agents that have a goal, and a cost
    - E.g., reach goal with lowest cost
  - Later: agents that have numerical utilities, rewards, etc.
    - E.g., take actions that maximize total reward over time (e.g., largest profit in \$)

# Agent design

---

- The environment type largely determines the agent design
  - *Fully/partially observable* => agent requires **memory** (internal state)
  - *Discrete/continuous* => agent may not be able to enumerate **all states**
  - *Stochastic/deterministic* => agent may have to prepare for **contingencies**
  - *Single-agent/multi-agent* => agent may need to behave **randomly**

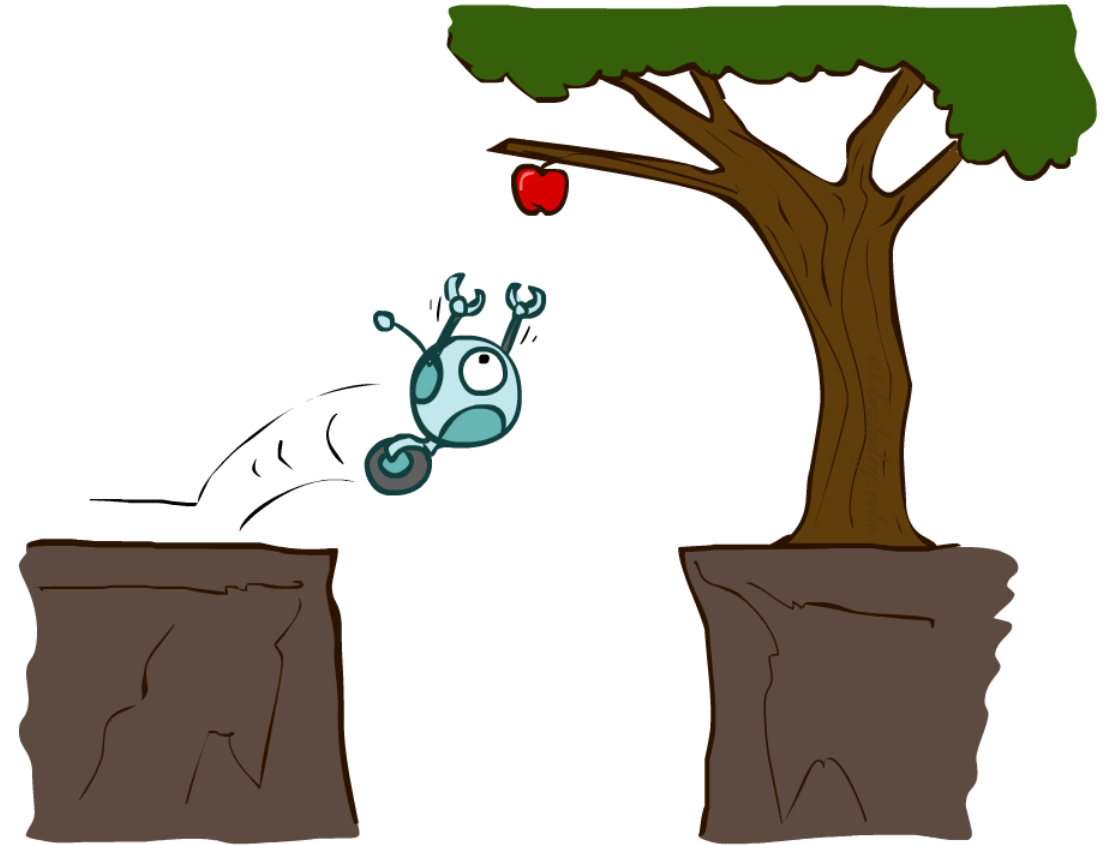
# Agents that Plan





# Reflex Agents

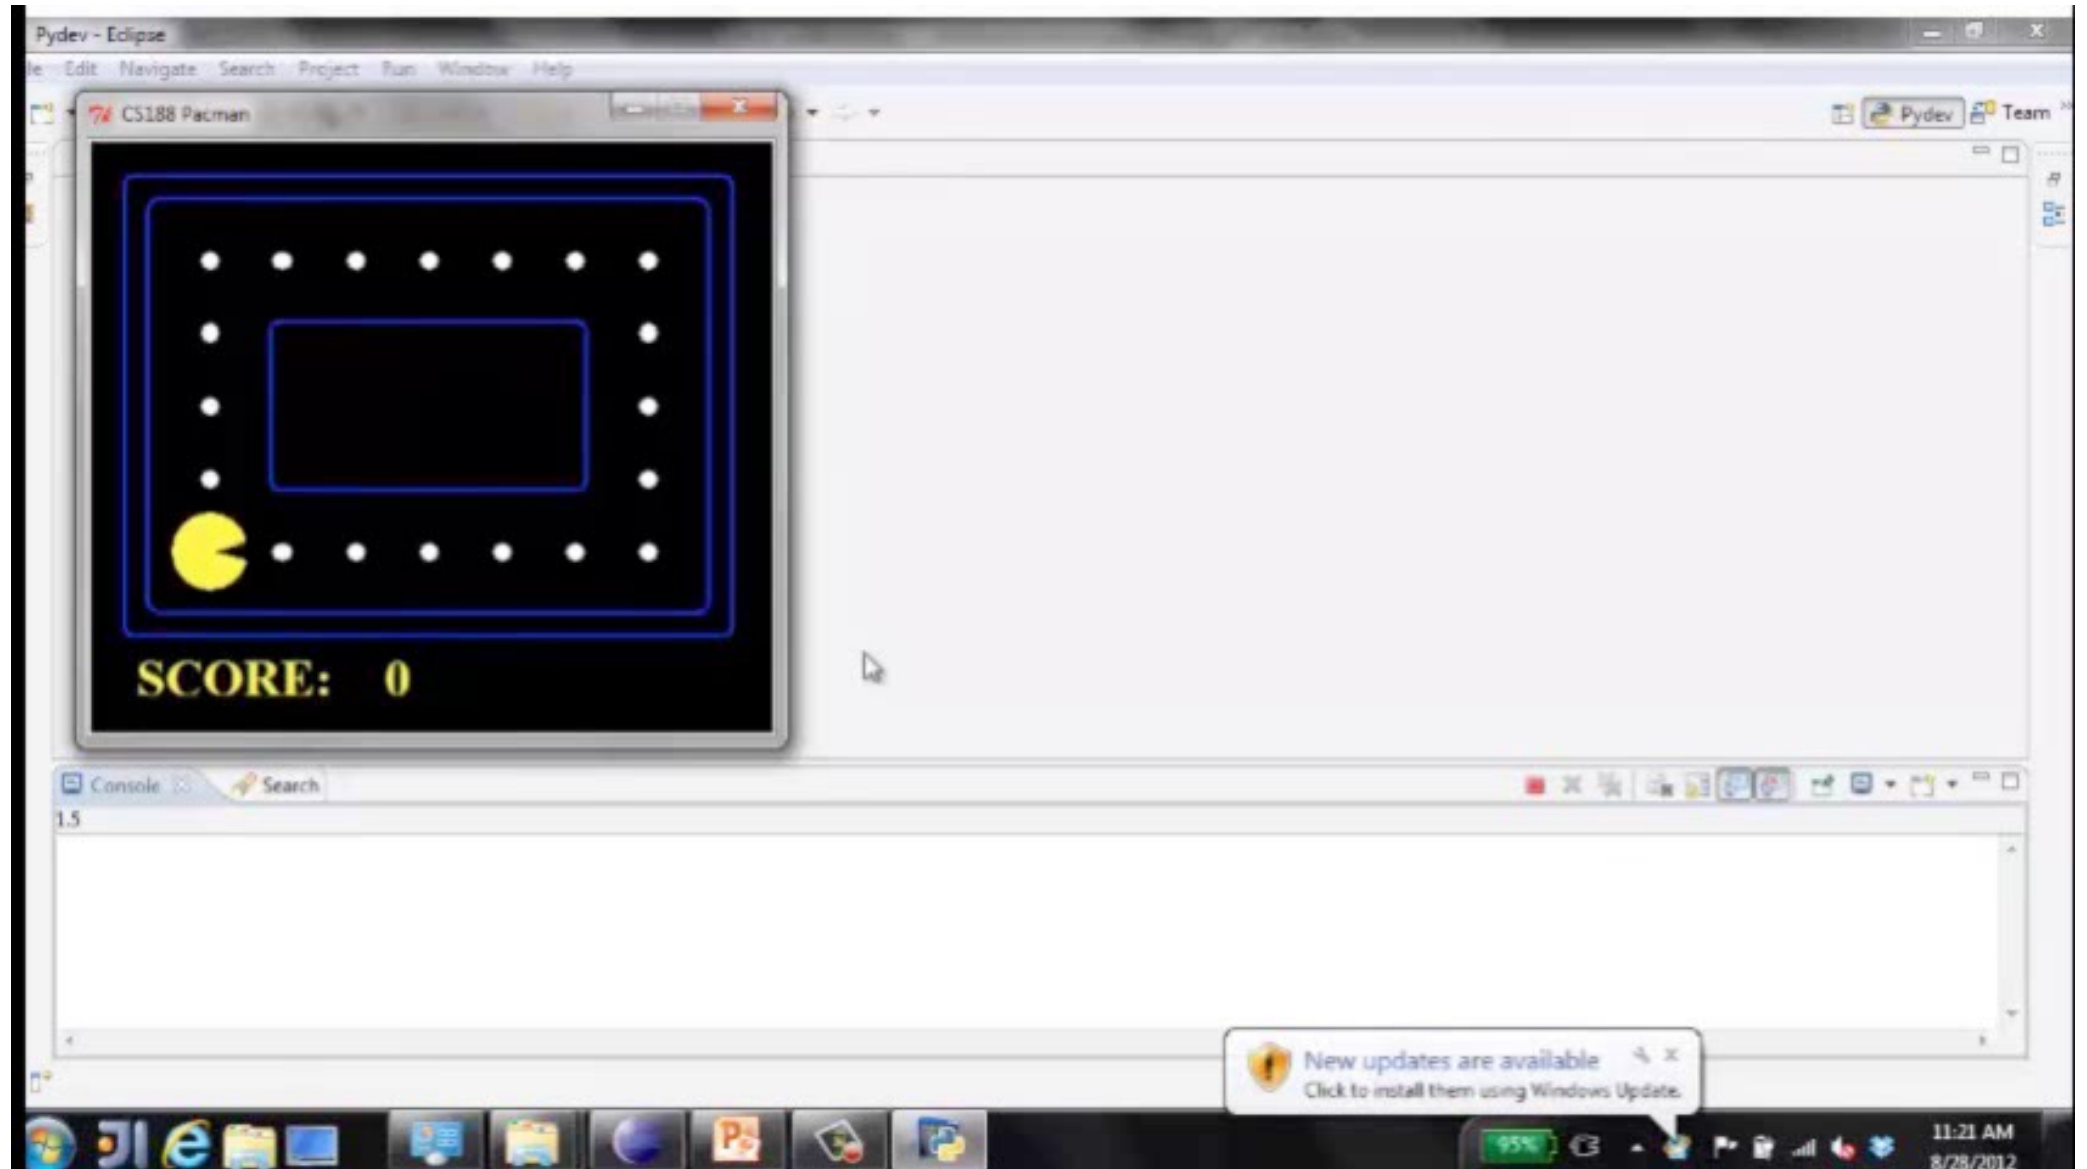
- Reflex agents:
  - Choose action based on current percept (and maybe memory)
  - May have memory or a model of the world's current state
  - Do not consider the future consequences of their actions
  - Consider how the world IS
- Can a reflex agent be rational?



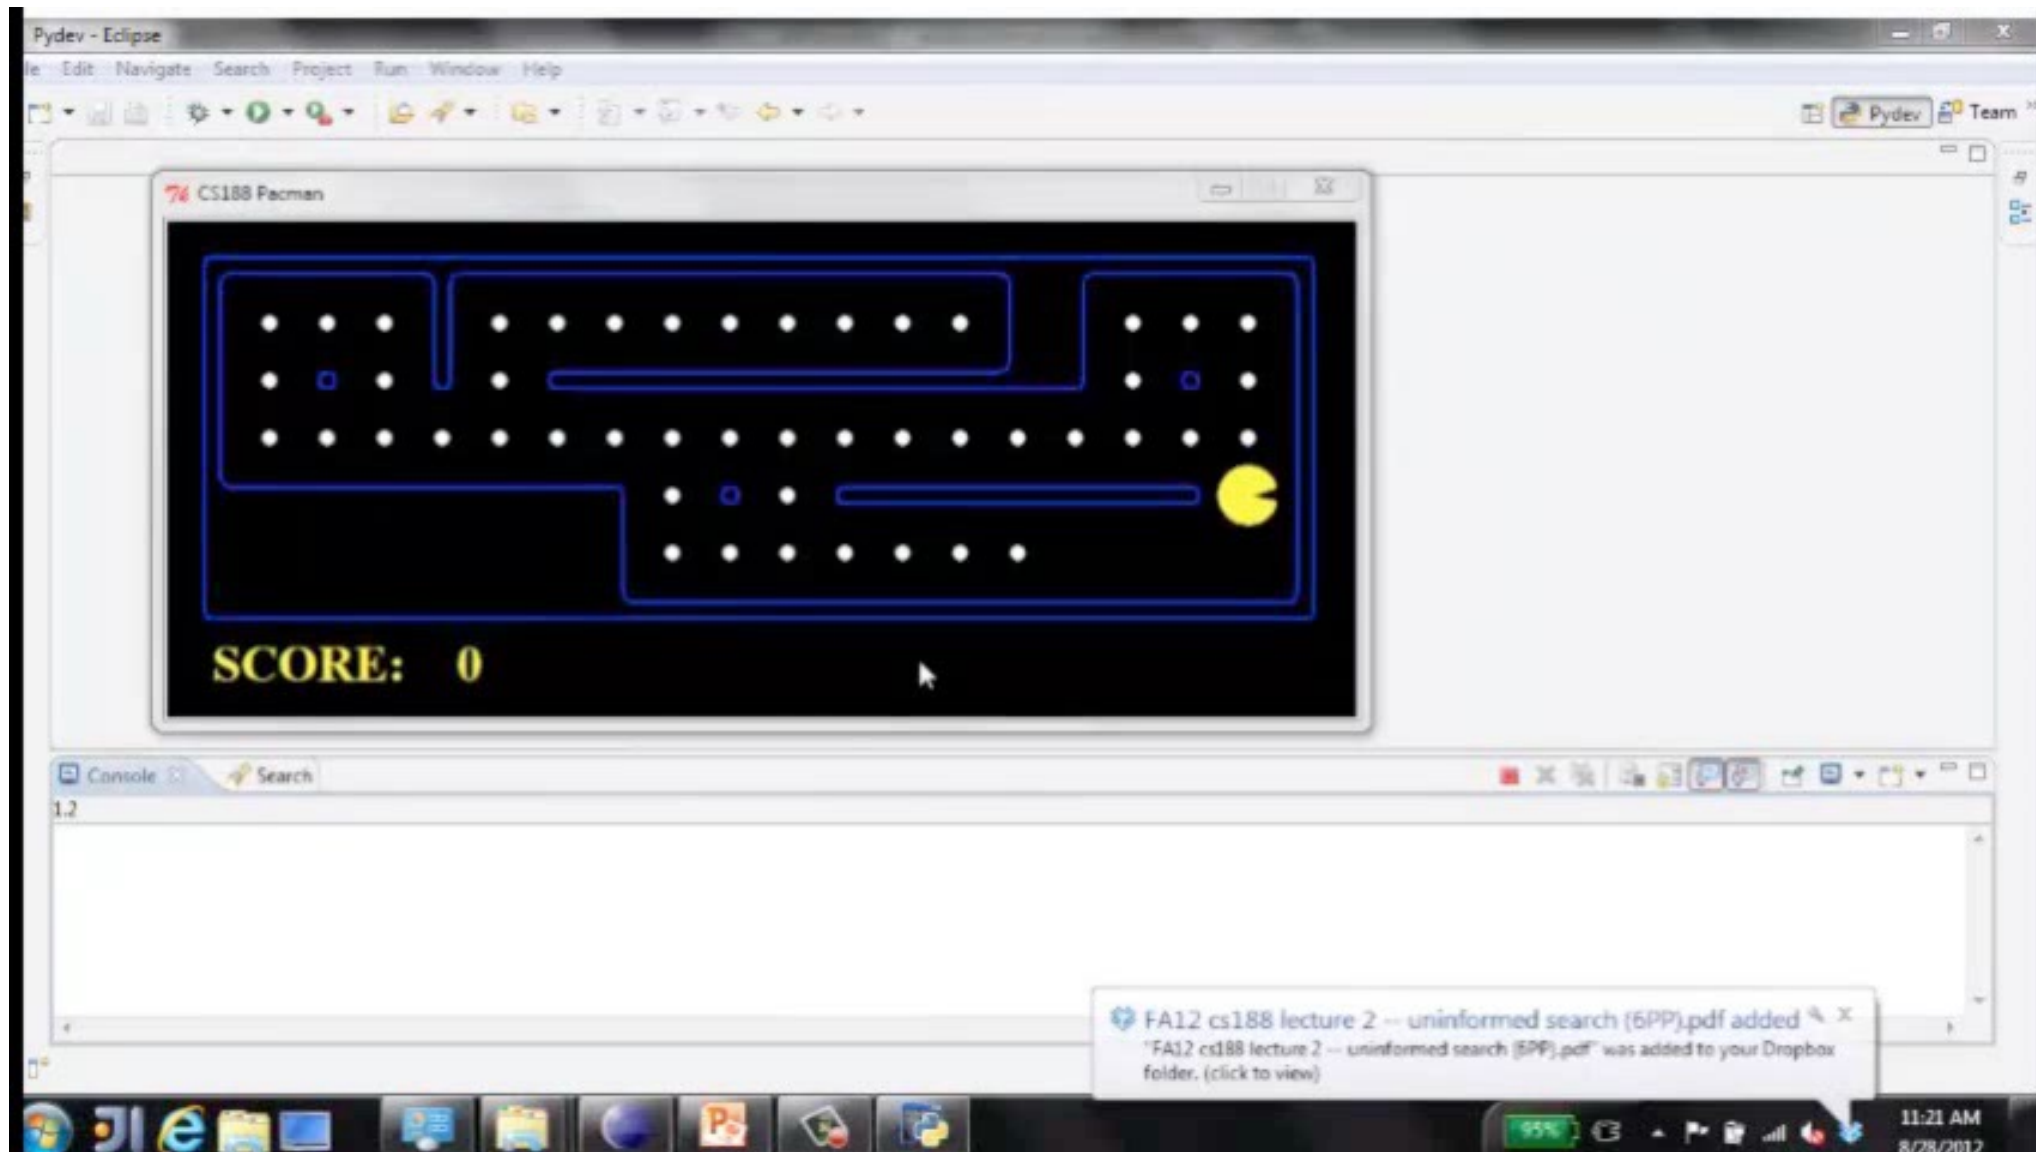
[Demo: reflex optimal (L2D1)]

[Demo: reflex optimal (L2D2)]

# Video of Demo Reflex Optimal



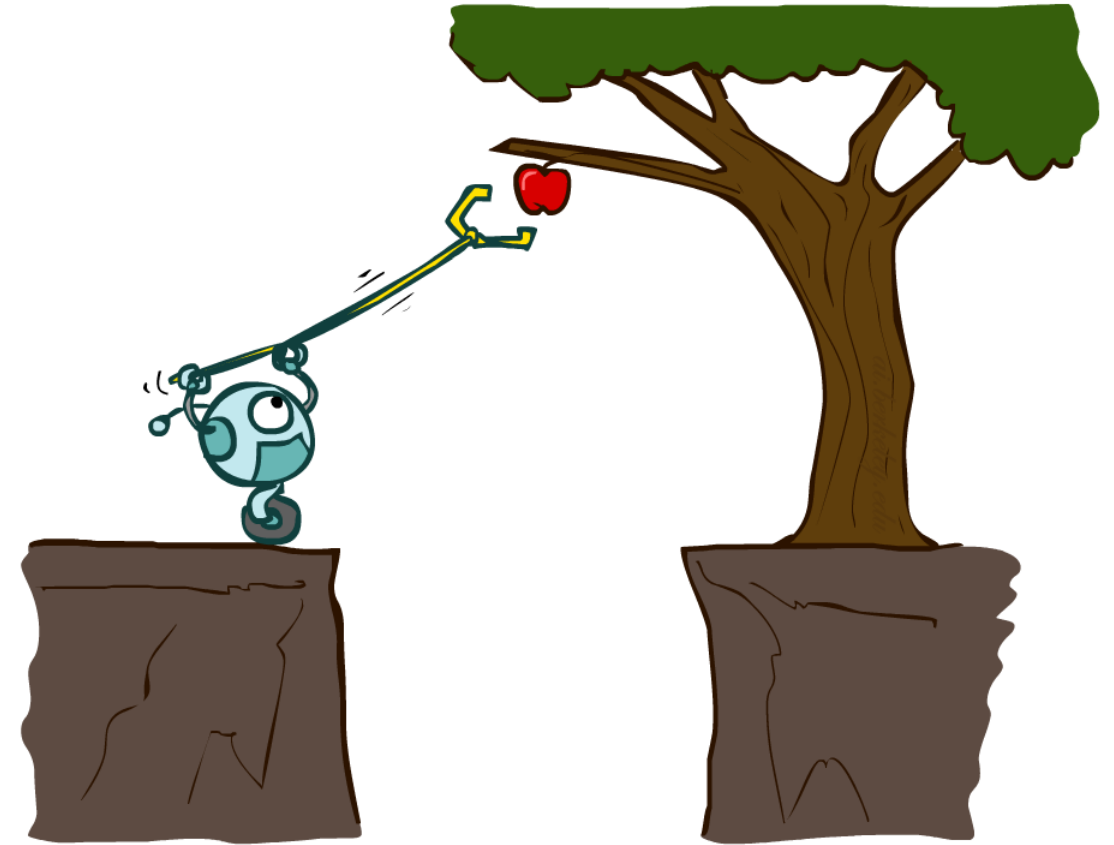
# Video of Demo Reflex Odd





# Planning Agents

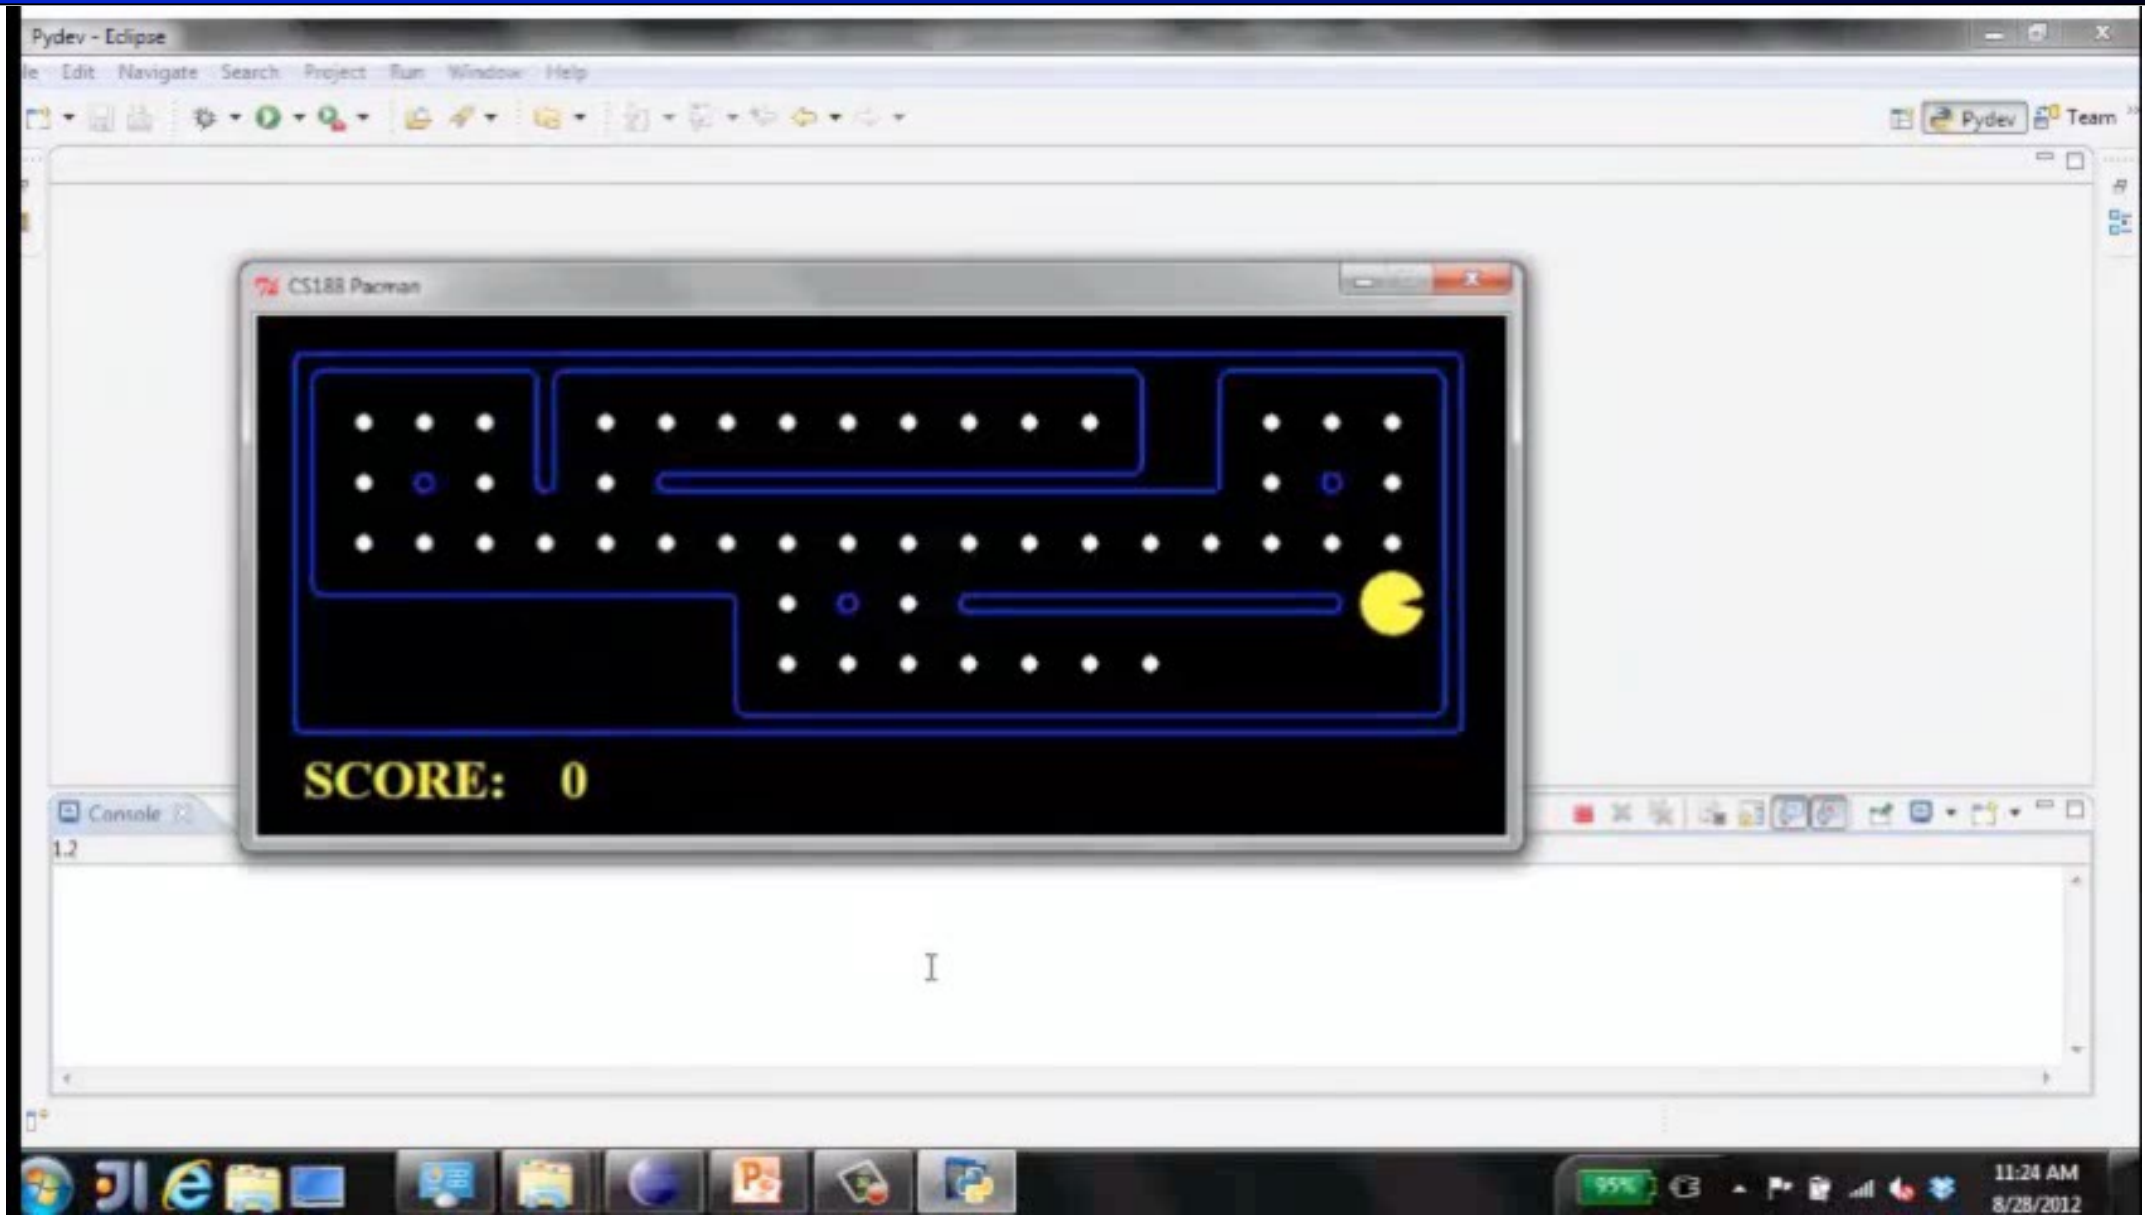
- Planning agents:
  - Ask “what if”
  - Decisions based on (hypothesized) consequences of actions
  - Must have a model of how the world evolves in response to actions
  - Must formulate a goal (test)
  - Consider how the world **WOULD BE**
- Optimal vs. complete planning
- Planning vs. replanning



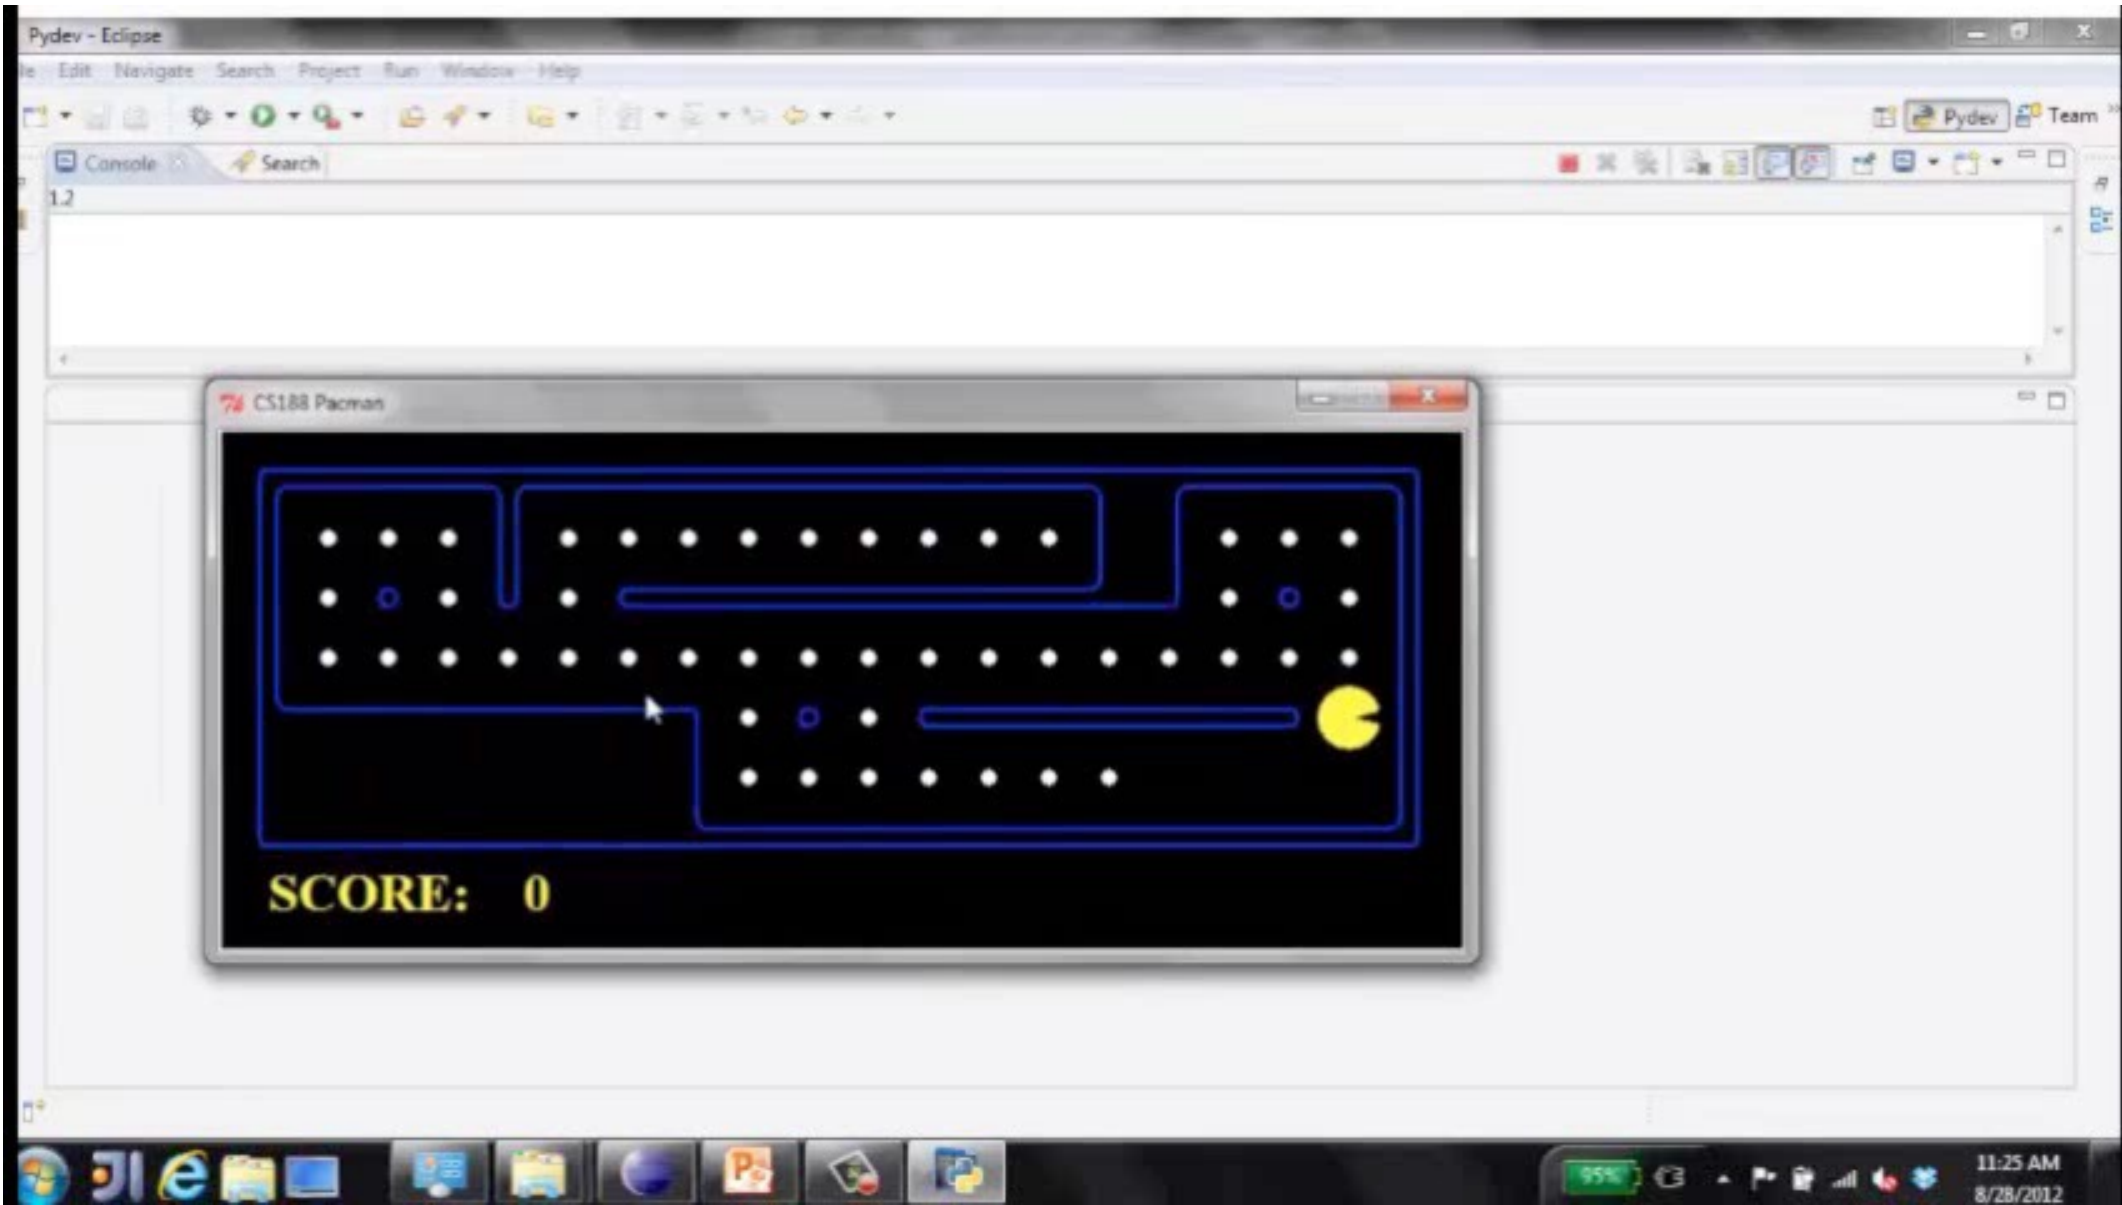
[Demo: re-planning (L2D3)]

[Demo: mastermind (L2D4)]

# Video of Demo Replanning



# Video of Demo Mastermind





# Search Problems

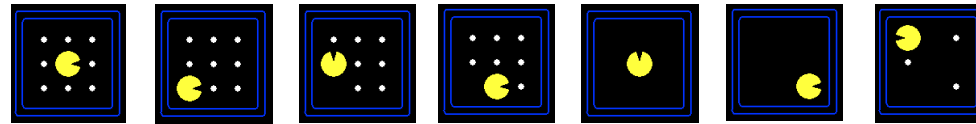
---



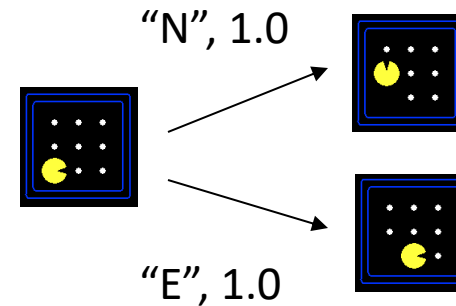
# Search Problems

- A **search problem** consists of:

- A state space



- A successor function  
(with actions, costs)



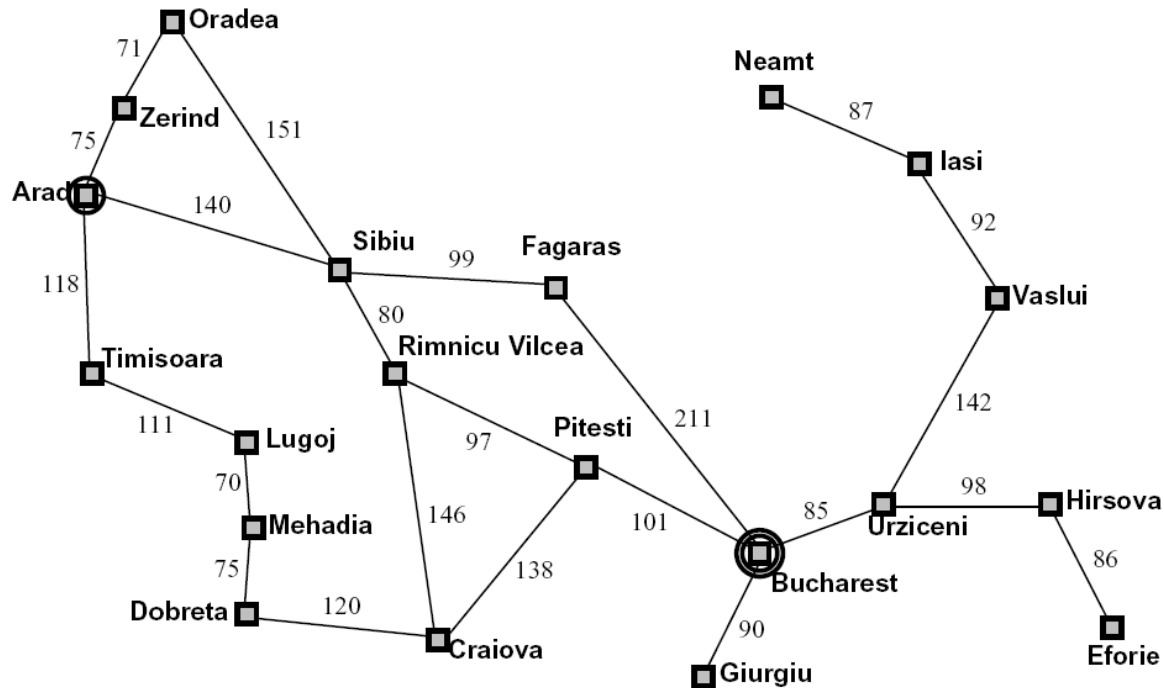
- A start state and a goal test
- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

# Search Problems Are Models

---



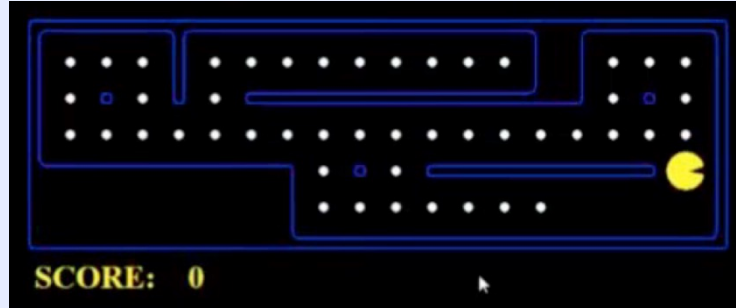
# Example: Traveling in Romania



- State space:
  - Cities
- Successor function:
  - Roads: Go to adjacent city with cost = distance
- Start state:
  - Arad
- Goal test:
  - Is state == Bucharest?
- Solution?

# What's in a State Space?

The **world state** includes every last detail of the environment



A **search state** keeps only the details needed for planning (abstraction)

## ■ Problem: Pathing

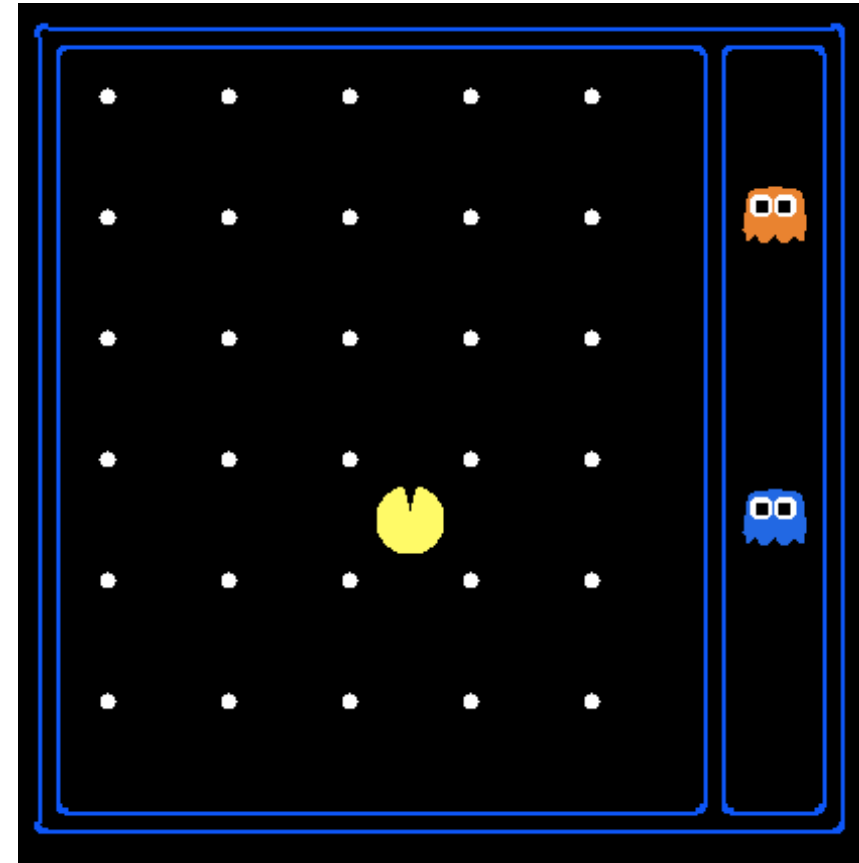
- States:  $(x,y)$  location
- Actions: NSEW
- Successor: update location only
- Goal test: is  $(x,y)=\text{END}$

## ■ Problem: Eat-All-Dots

- States:  $\{(x,y), \text{dot booleans}\}$
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false

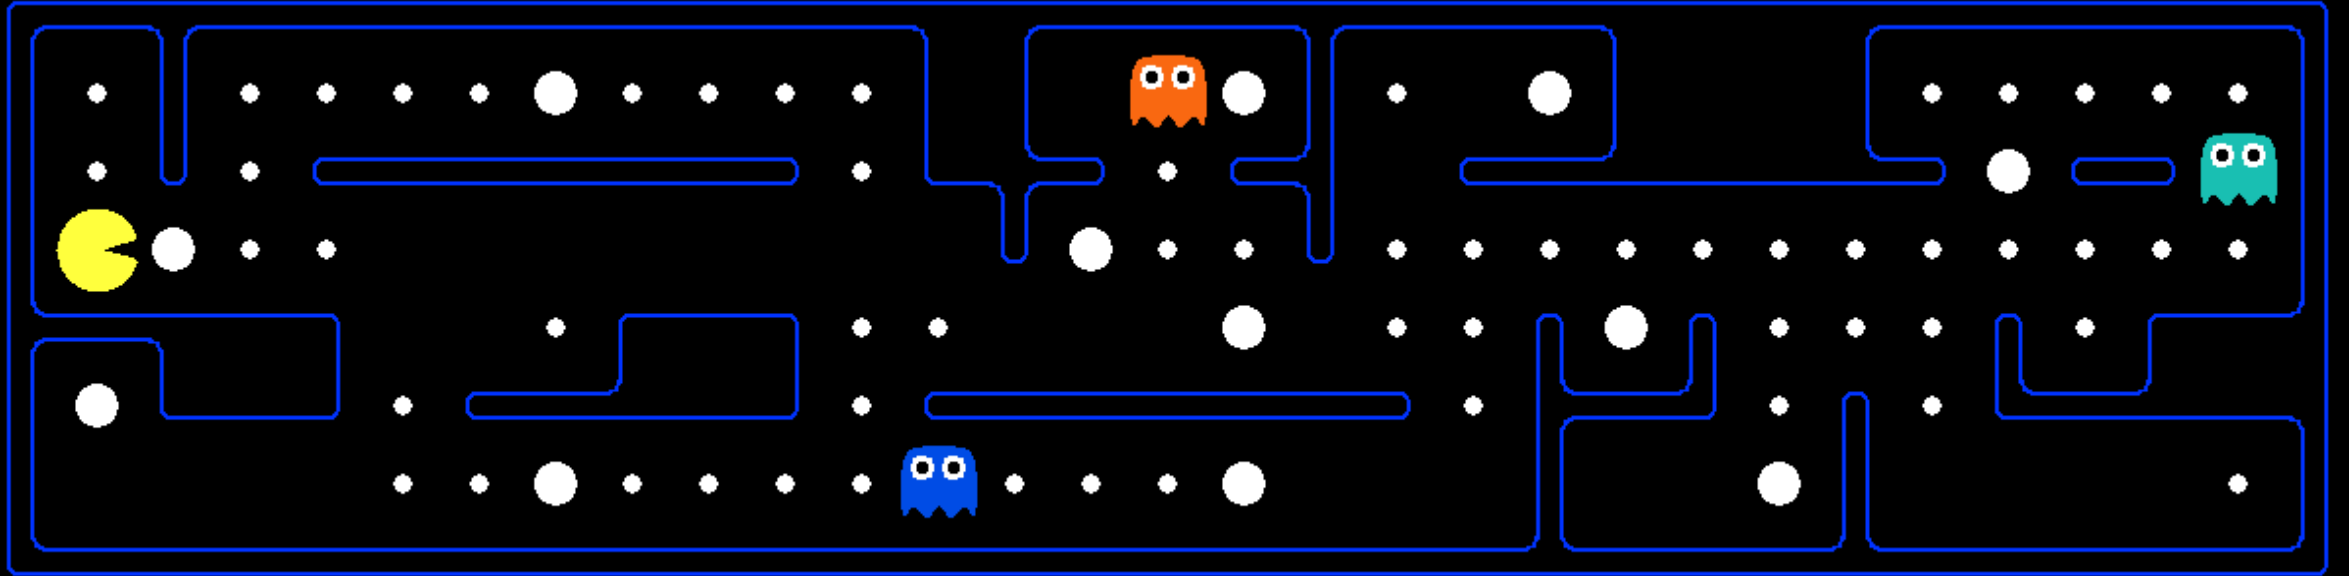
# State Space Sizes?

- World state:
  - Agent positions: 120
  - Food count: 30
  - Ghost positions: 12
  - Agent facing: NSEW
- How many
  - World states?  
 $120 \times (2^{30}) \times (12^2) \times 4$
  - States for pathing?  
120
  - States for eat-all-dots?  
 $120 \times (2^{30})$



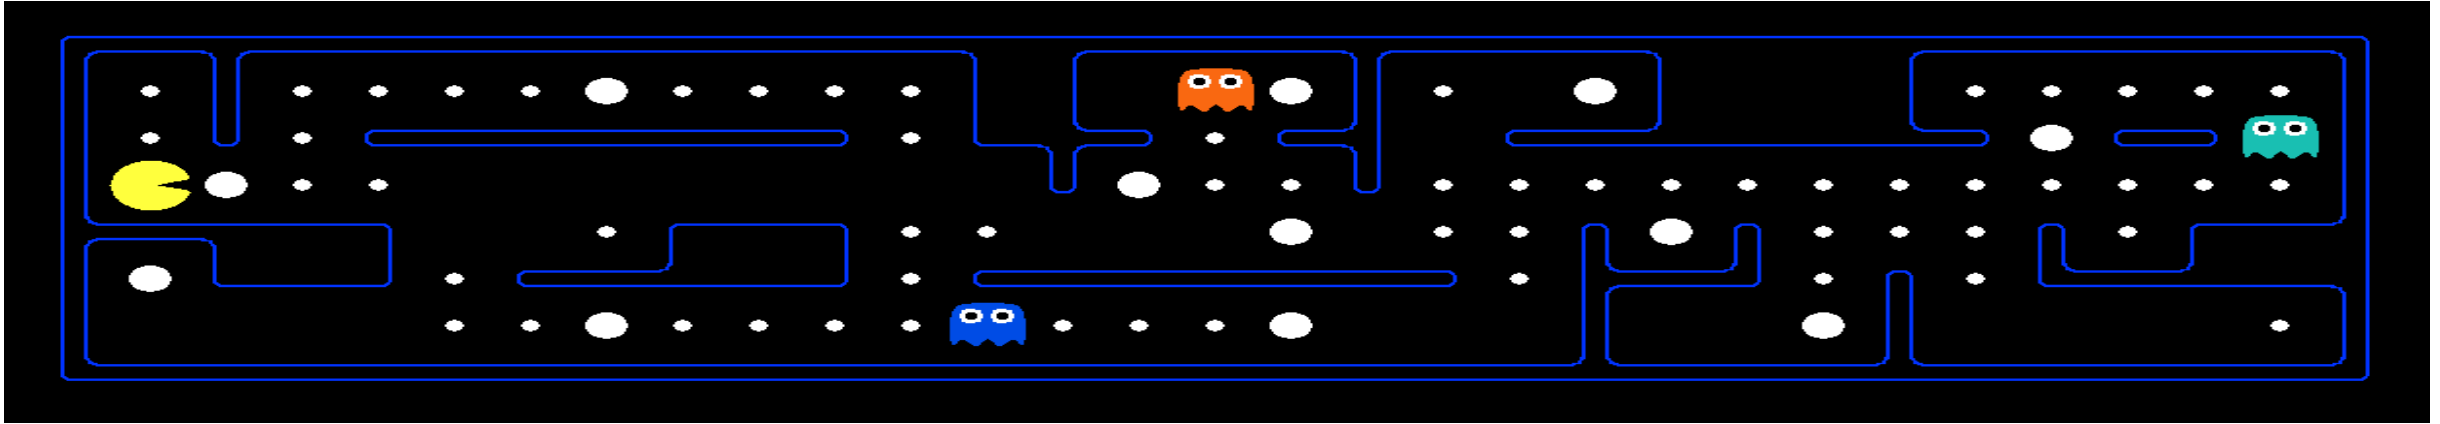


# Quiz: Safe Passage



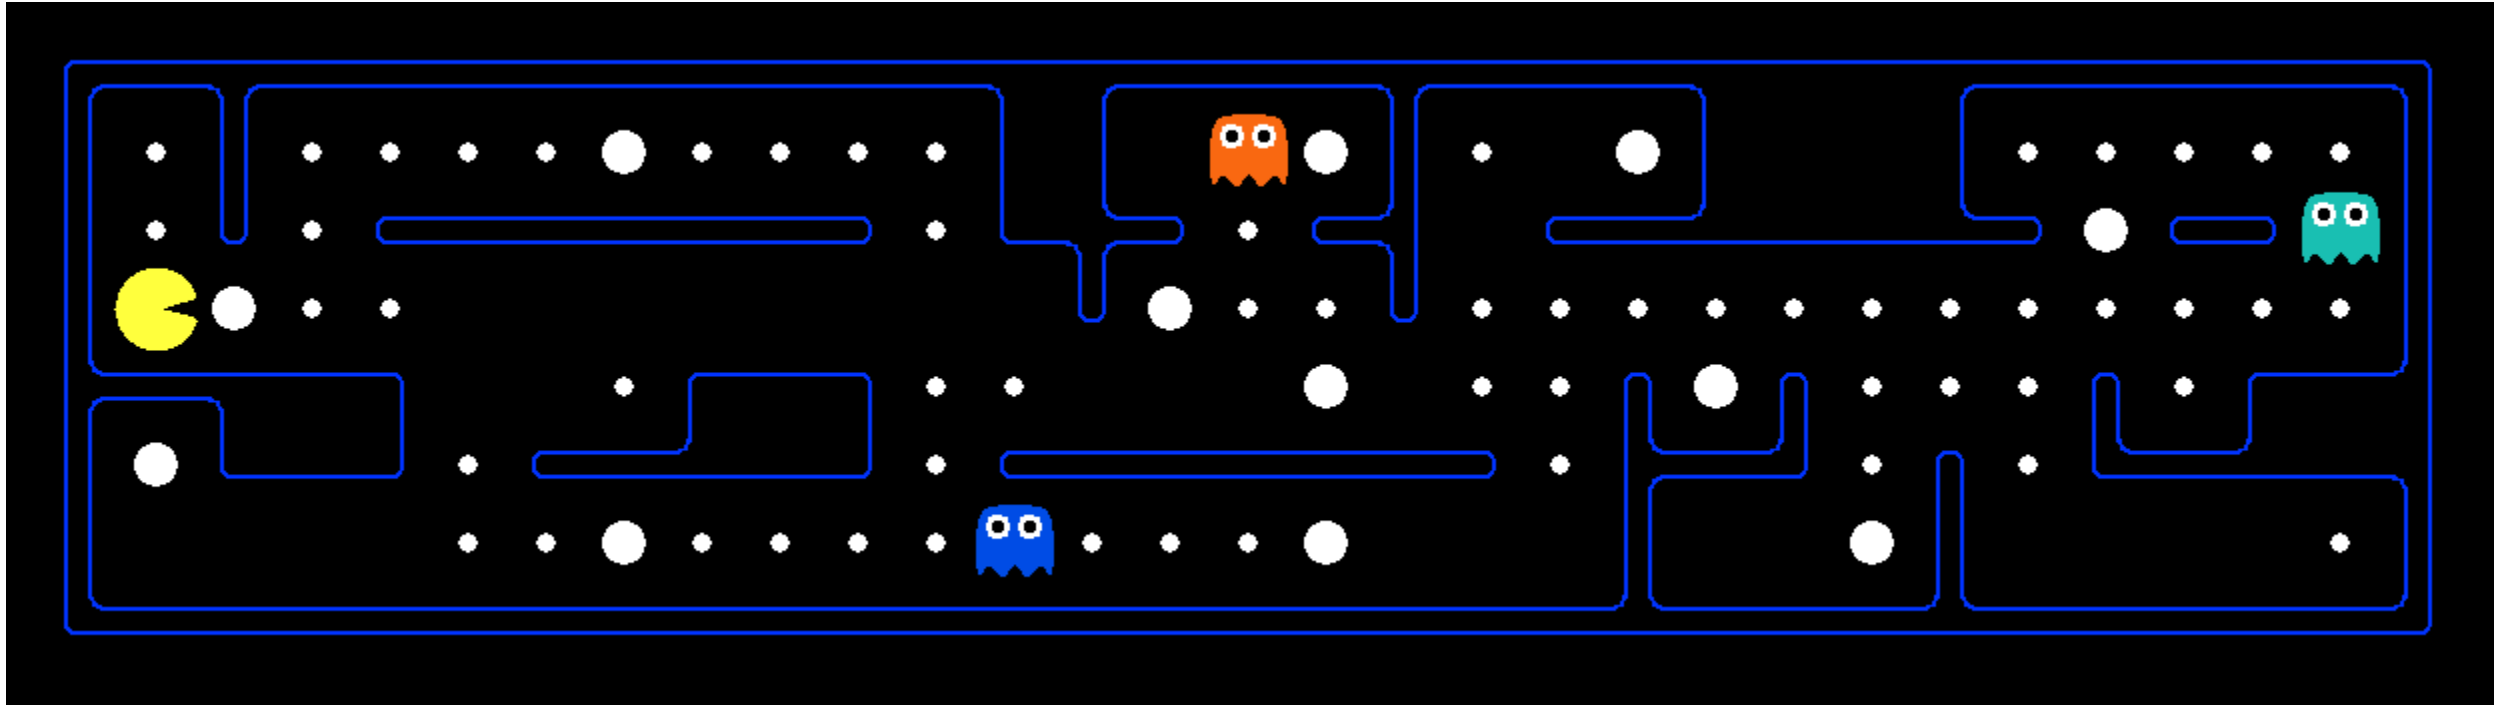
- Problem: eat all dots while keeping the ghosts perma-scared
- What does the state space have to specify?

# Options



- Problem: eat all dots while keeping the ghosts perma-scared
- What does the state space have to specify? Pick from the list:
  1. Pacman's position
  2. Postion of each ghost
  3. A boolean for each dot that represents whether it has been eaten
  4. A boolean for each power pellet that represents whether it has been eaten
  5. The remaining amount of time for which the ghosts will be

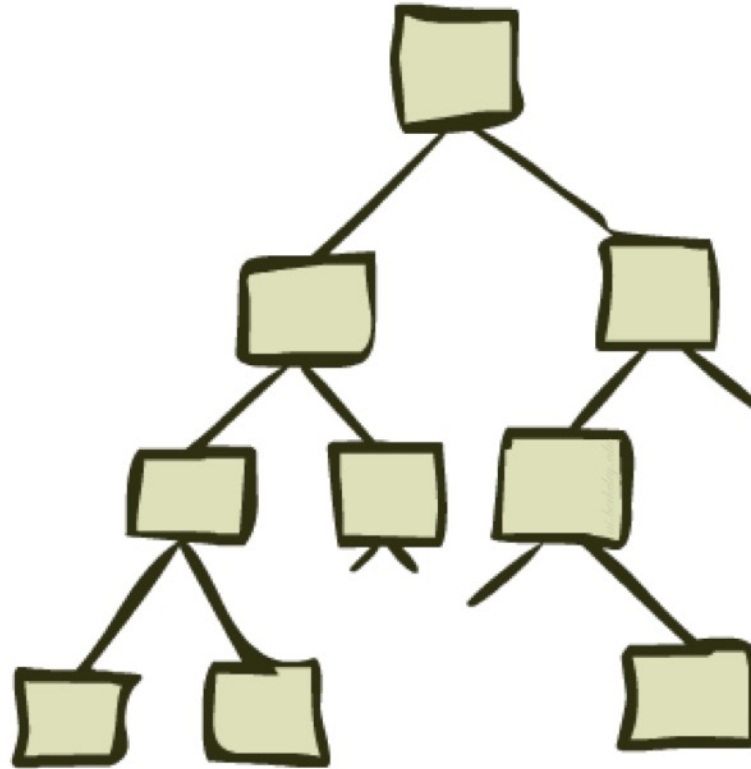
# Quiz: Safe Passage



- Problem: eat all dots while keeping the ghosts perma-scared
- What does the state space have to specify?
  - (agent position, dot booleans, power pellet booleans, remaining scared time)

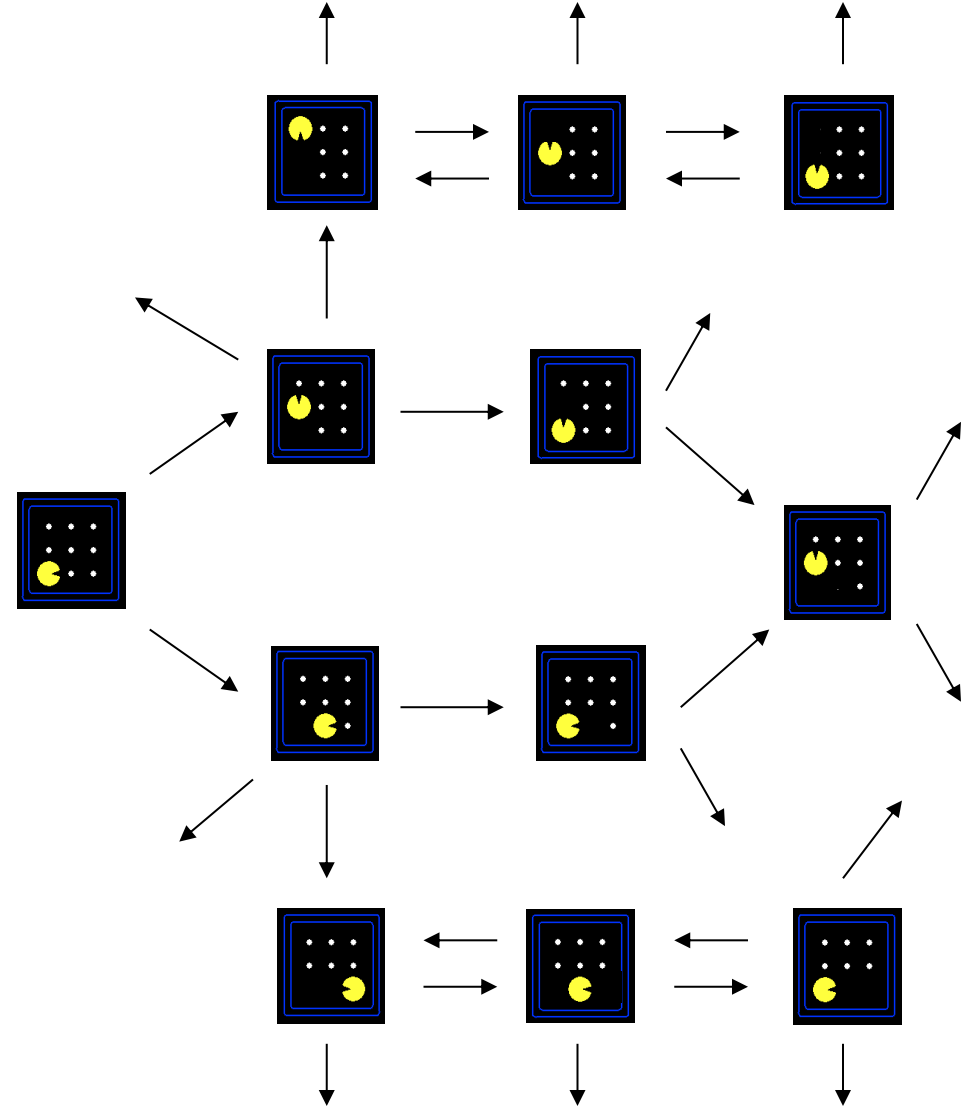
# State Space Graphs and Search Trees

---



# State Space Graphs

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



# Legend of Chessboard

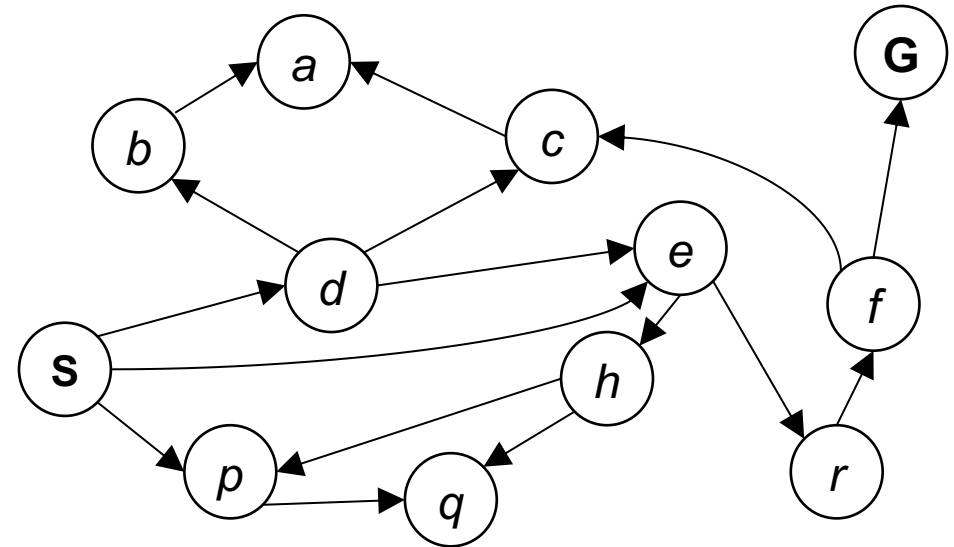
---





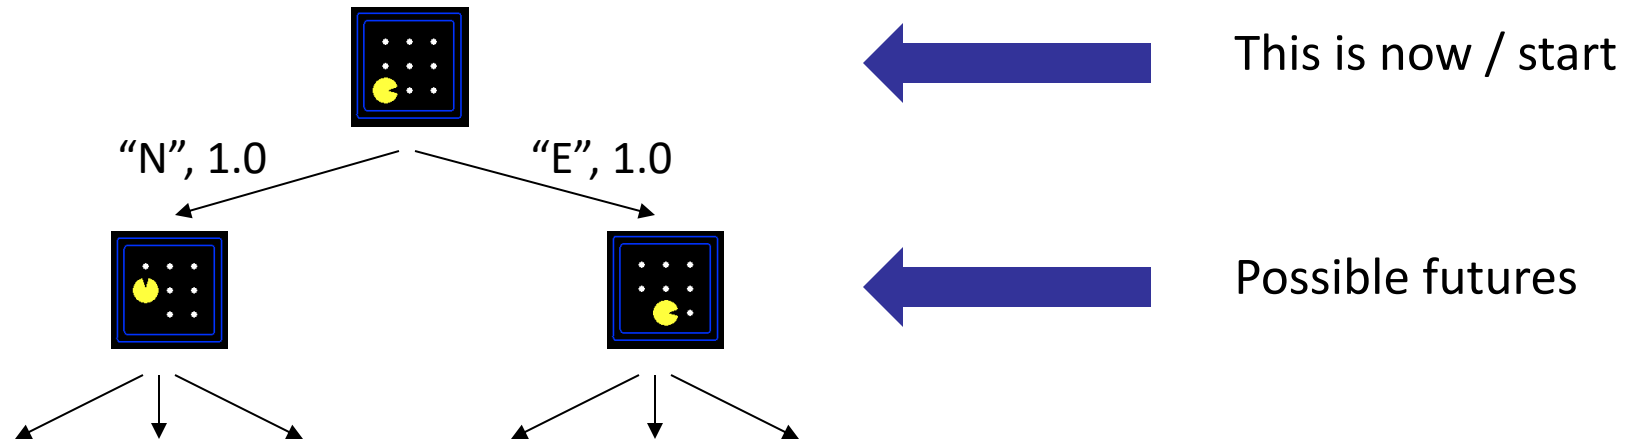
# State Space Graphs

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



*Tiny state space graph for a tiny search problem*

# Search Trees

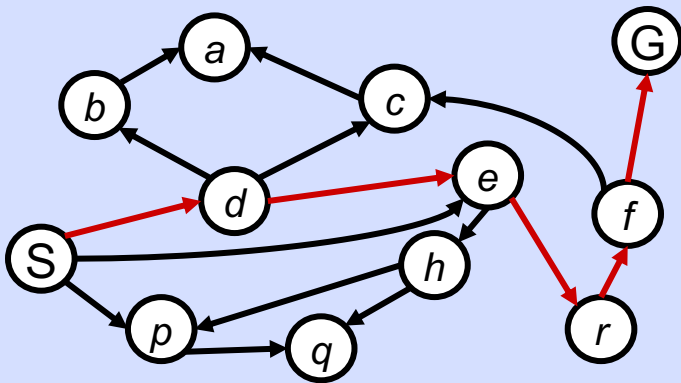


- A search tree:

- A “what if” tree of plans and their outcomes
- The start state is the root node
- Children correspond to successors
- Nodes show states, but correspond to PLANS that achieve those states
- For most problems, we can never actually build the whole tree

# State Space Graphs vs. Search Trees

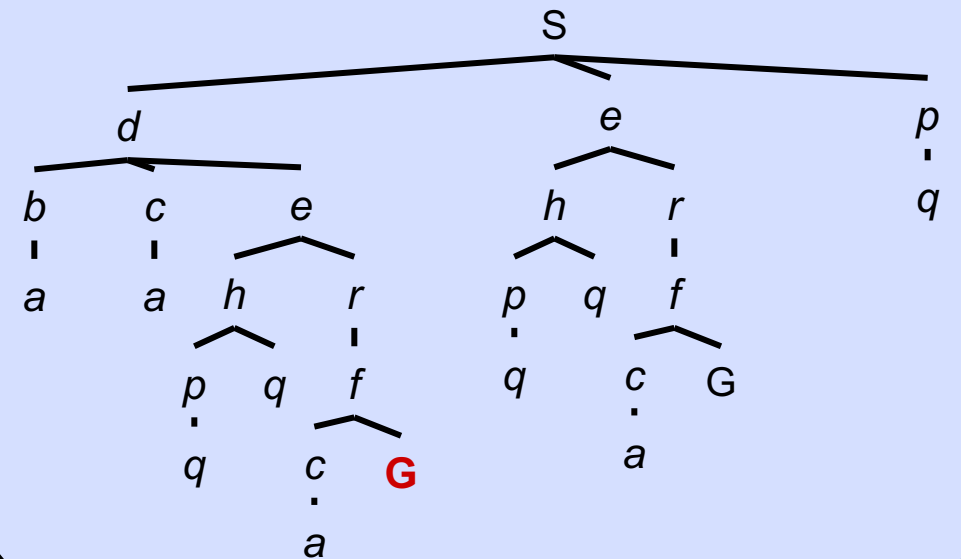
# State Space Graph



*Each NODE in in  
the search tree is  
an entire PATH in  
the state space  
graph.*

*We construct both  
on demand – and  
we construct as  
little as possible.*

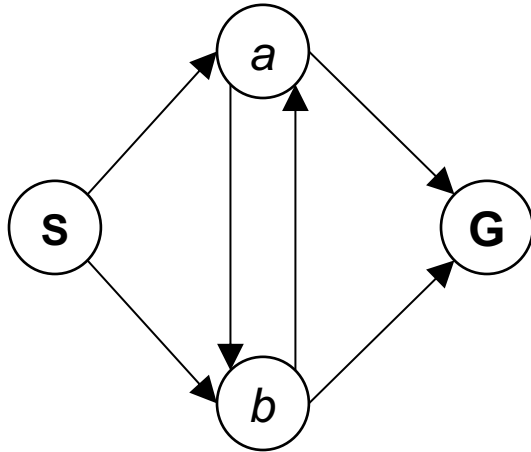
# Search Tree



# Quiz: State Space Graphs vs. Search Trees

---

Consider this 4-state graph:

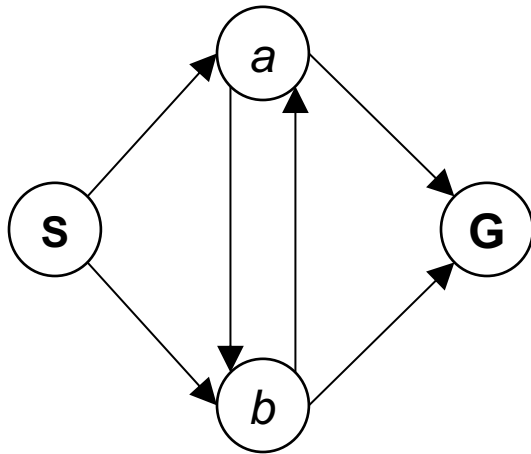


How big is its search tree (from S)?

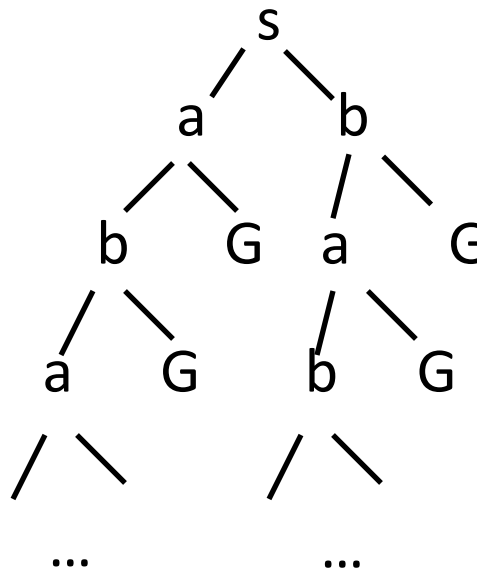


# Quiz: State Space Graphs vs. Search Trees

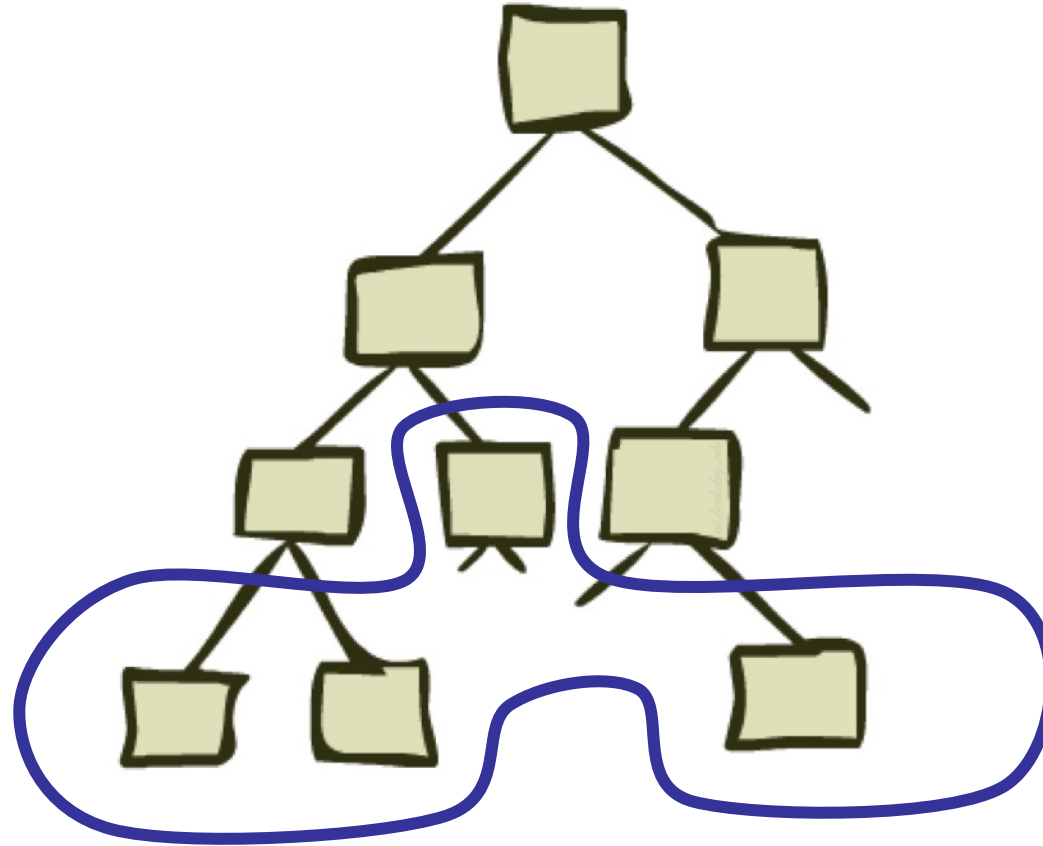
Consider this 4-state graph:



How big is its search tree (from S)?

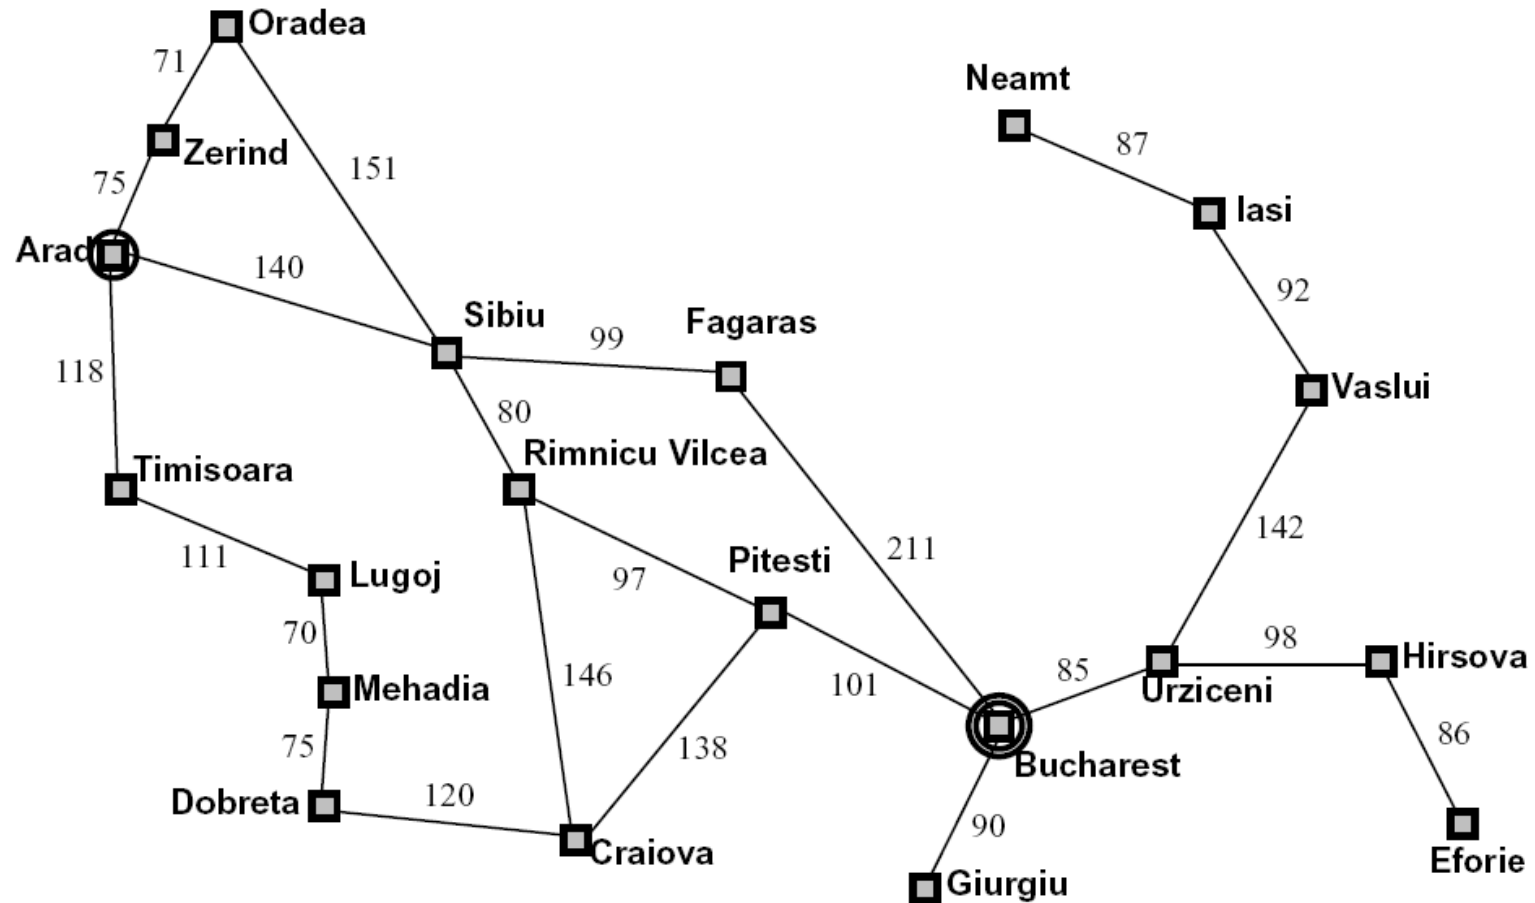


Important: Lots of repeated structure in the search tree!

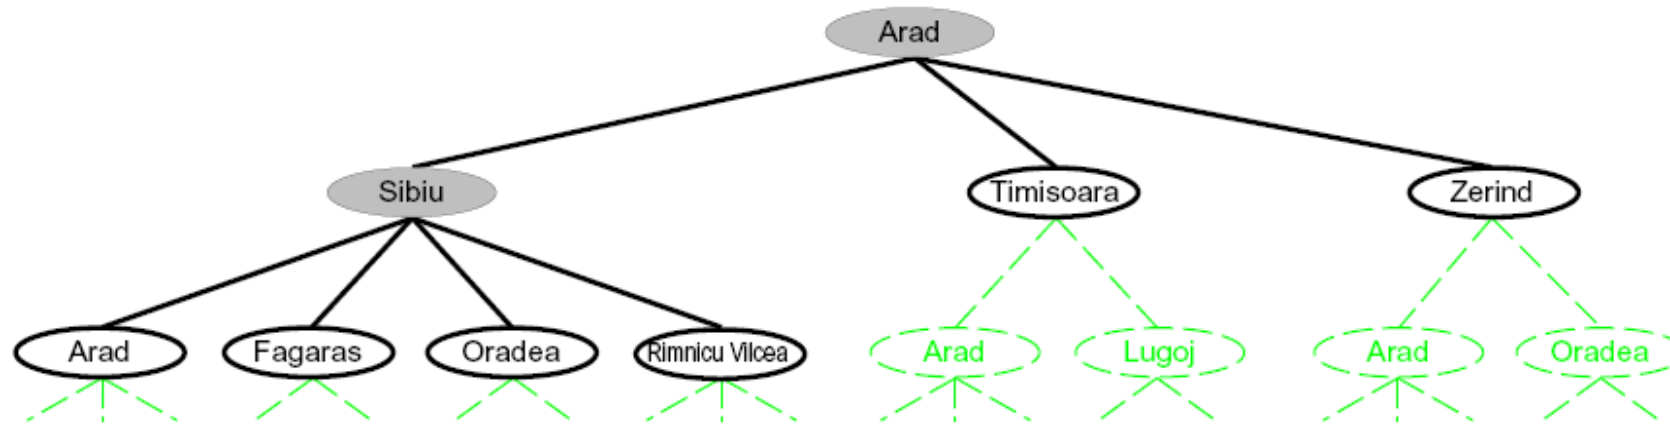




# Search Example: Romania



# Searching with a Search Tree



## ■ Search:

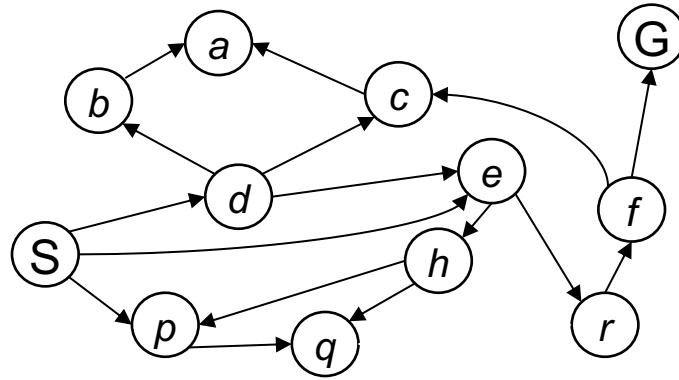
- Expand out potential plans (tree nodes)
- Maintain a **fringe** of partial plans under consideration
- Try to expand as few tree nodes as possible

# General Tree Search

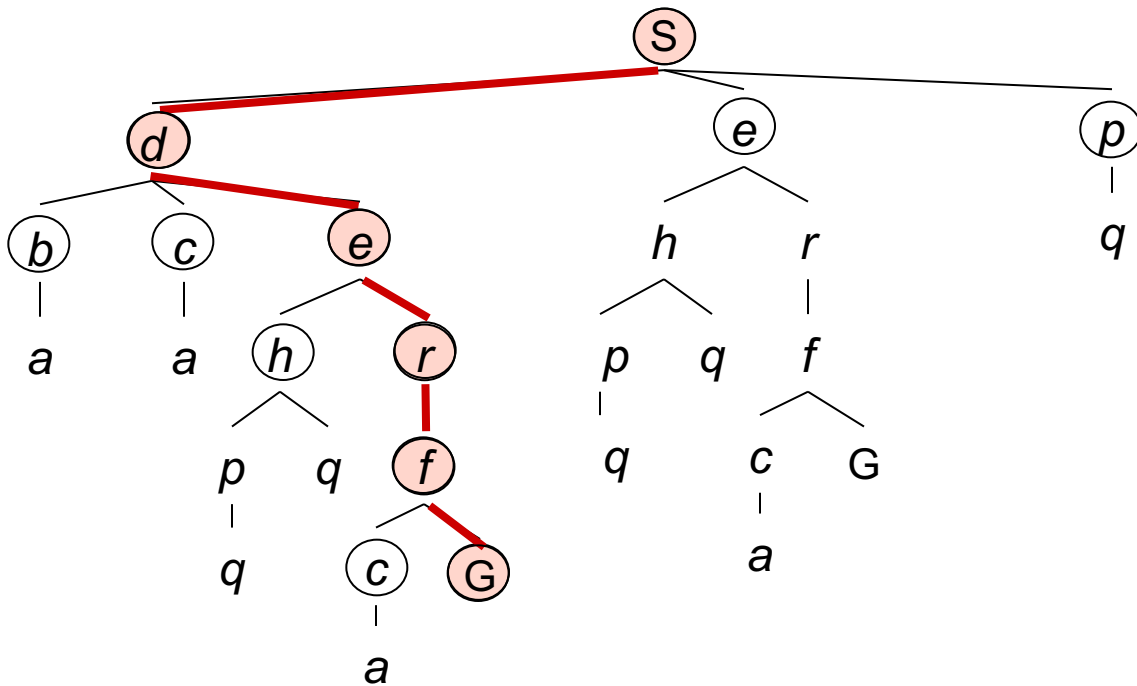
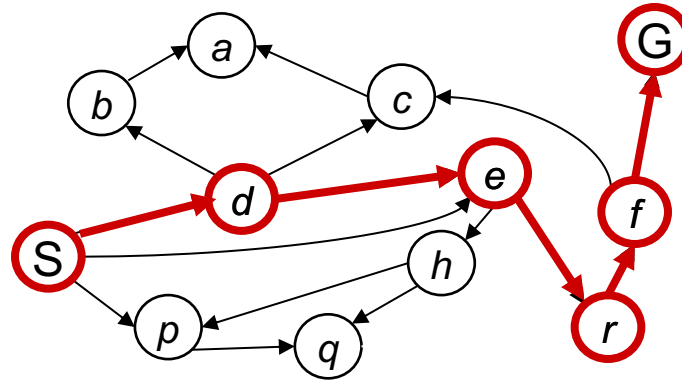
```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
  - Fringe
  - Expansion
  - Exploration strategy
- Main question: which fringe nodes to explore?

# Example: Tree Search



# Example: Tree Search



~~s~~  
~~s → d~~  
s → e  
s → p  
s → d → b  
s → d → c  
~~s → d → e~~  
s → d → e → h  
~~s → d → e → r~~  
~~s → d → e → r → f~~  
s → d → e → r → f → c  
~~s → d → e → r → f → G~~

# Depth-First Search

---

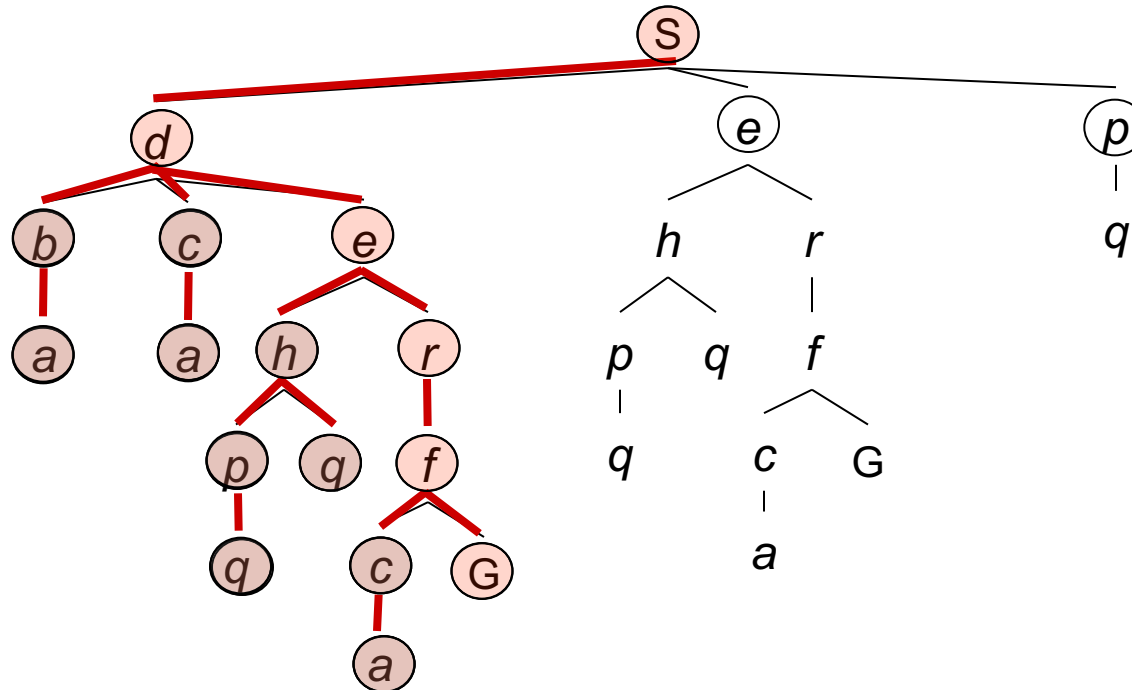
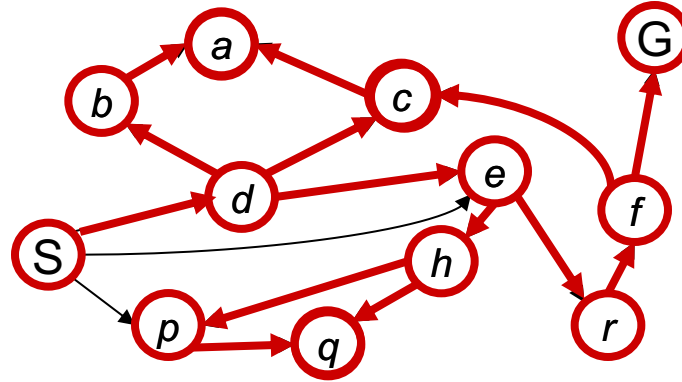




# Depth-First Search

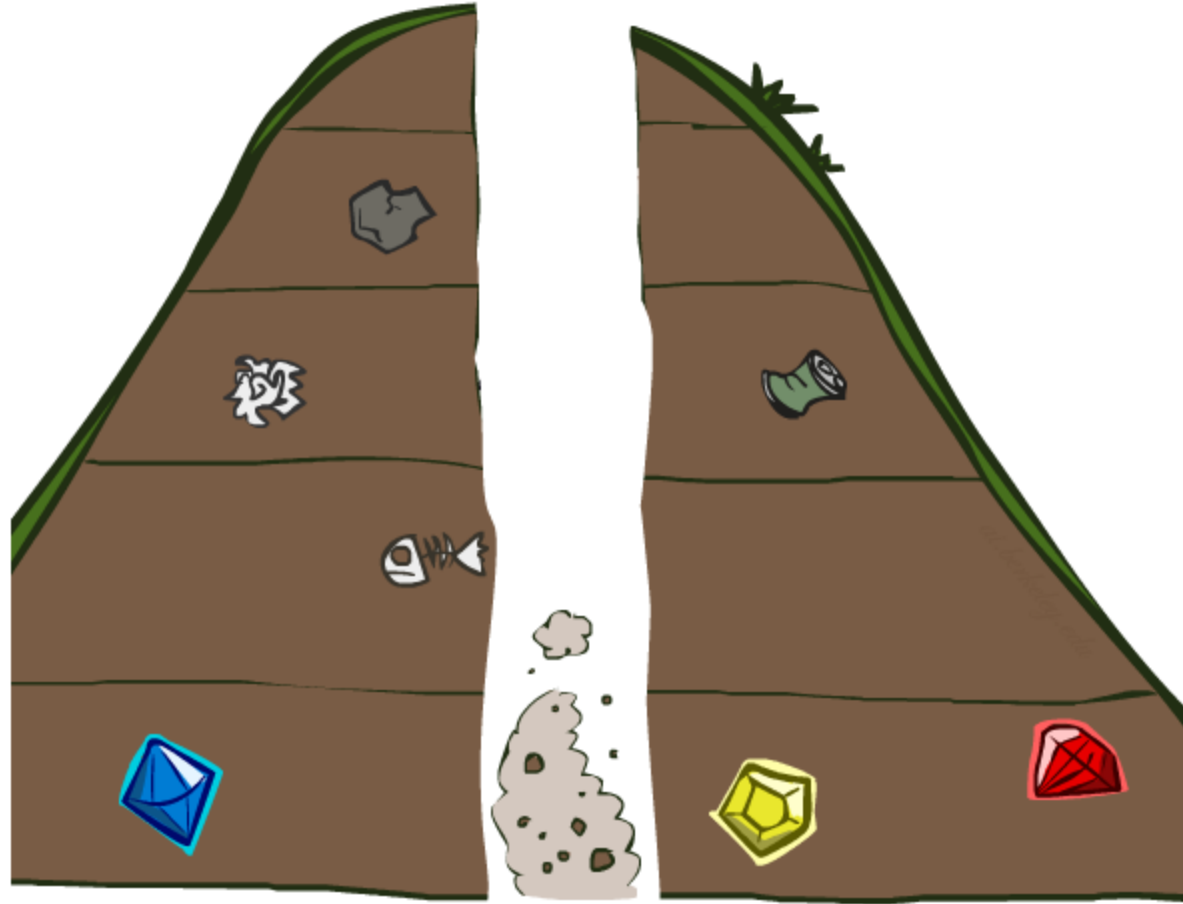
*Strategy: expand a  
deepest node first*

*Implementation:  
Fringe is a LIFO stack*



# Search Algorithm Properties

---



# Search Algorithm Properties

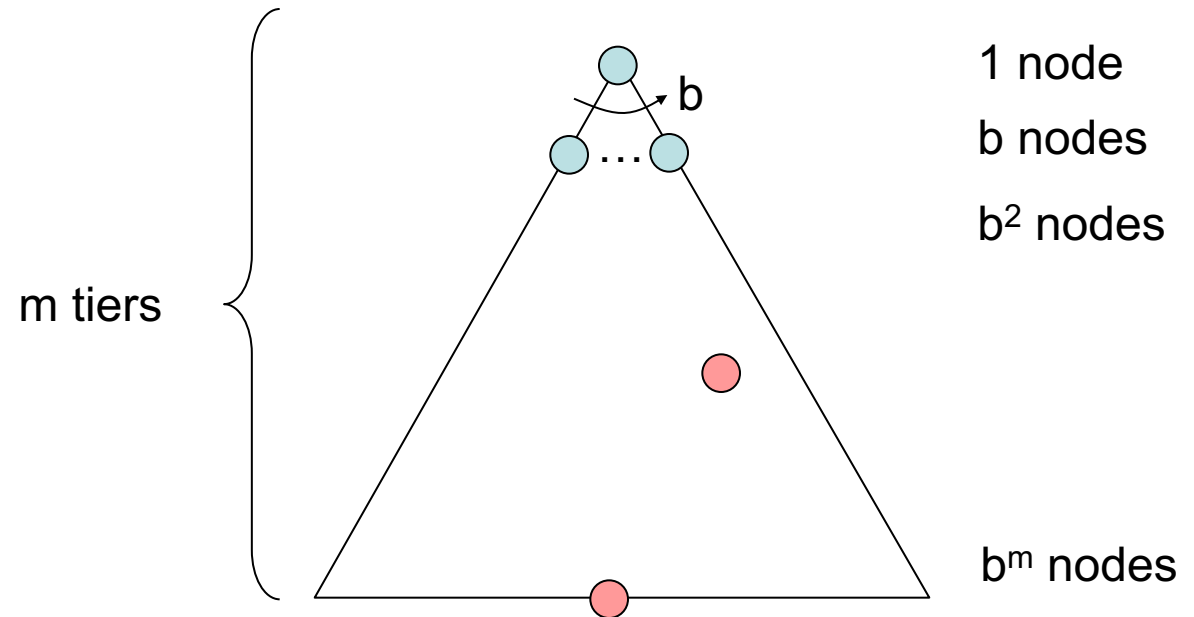
- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?

- Cartoon of search tree:

- $b$  is the branching factor
- $m$  is the maximum depth
- solutions at various depths

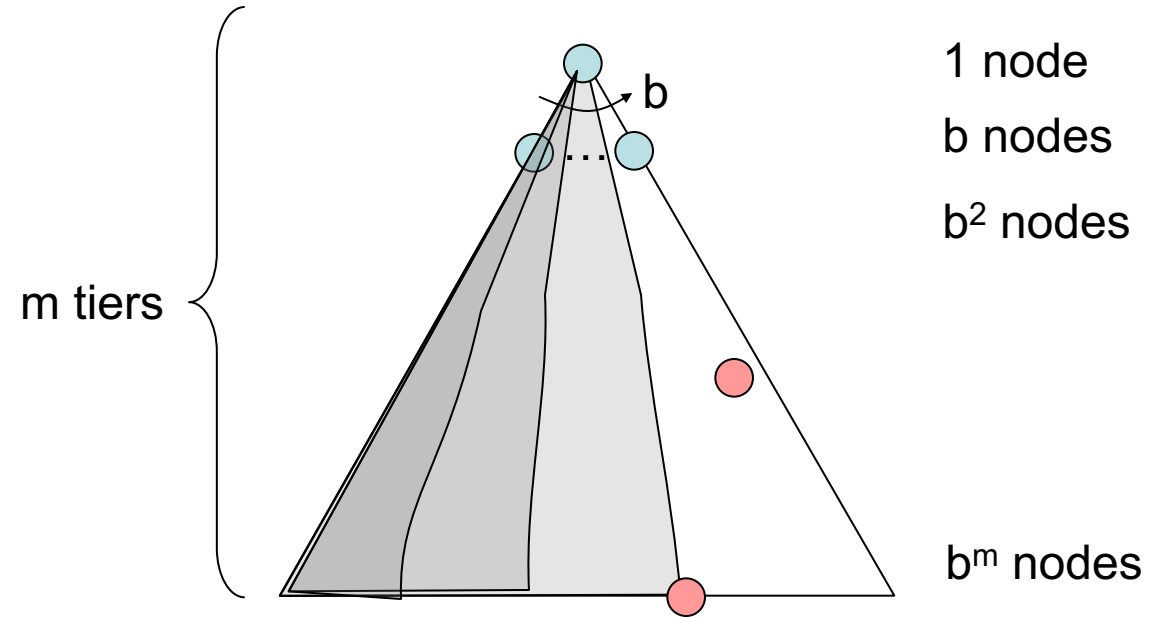
- Number of nodes in entire tree?

- $1 + b + b^2 + \dots + b^m = O(b^m)$



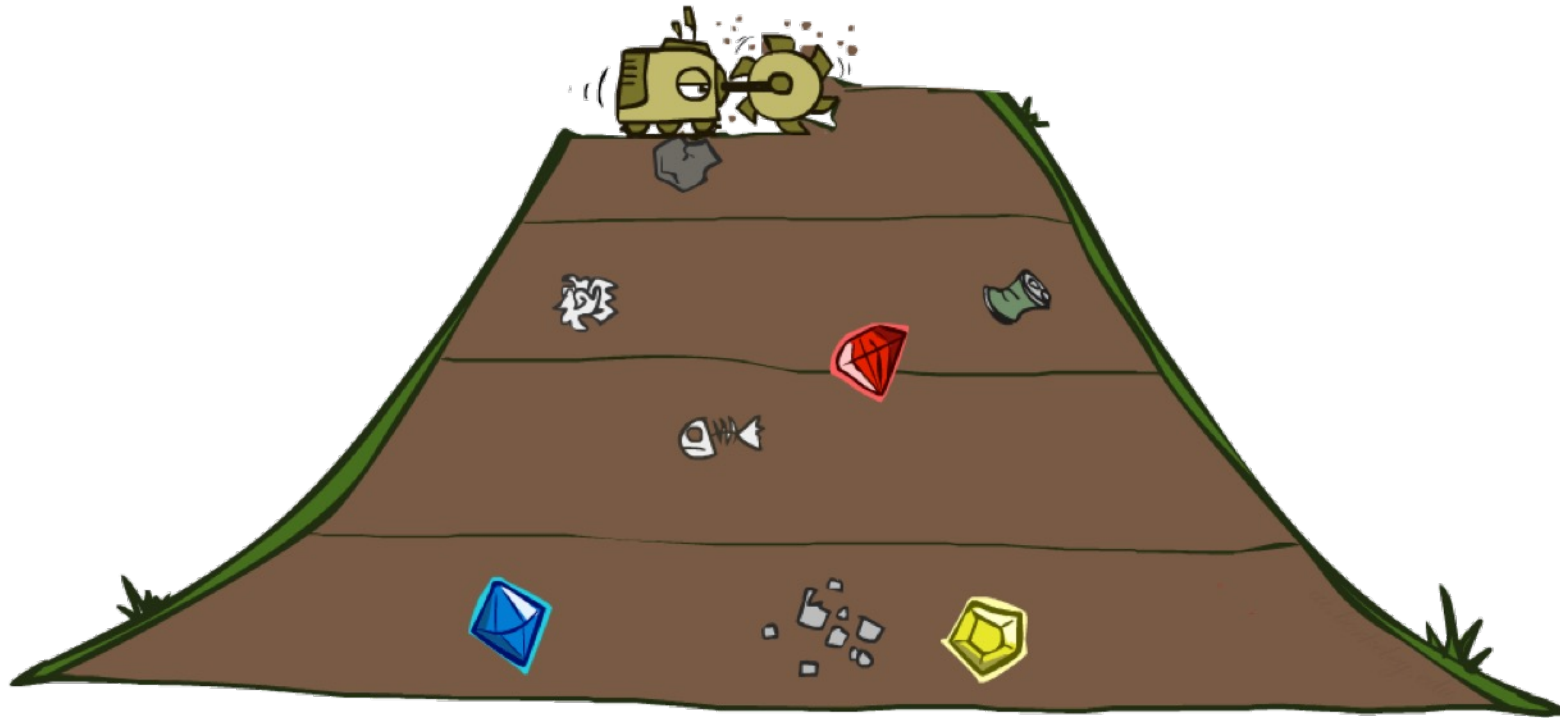
# Depth-First Search (DFS) Properties

- What nodes DFS expand?
  - Some left prefix of the tree.
  - Could process the whole tree!
  - If  $m$  is finite, takes time  $O(b^m)$
- How much space does the fringe take?
  - Only has siblings on path to root, so  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles (more later)
- Is it optimal?
  - No, it finds the “leftmost” solution, regardless of depth or cost



# Breadth-First Search

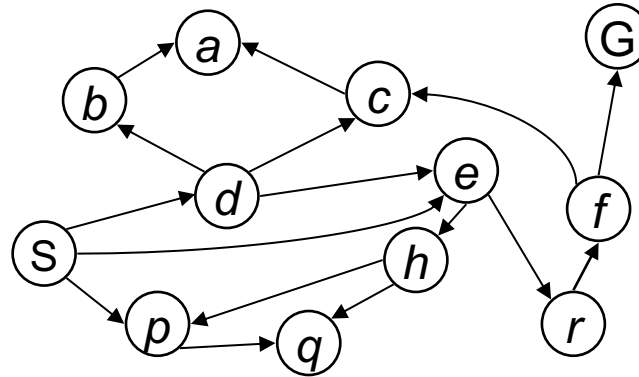
---



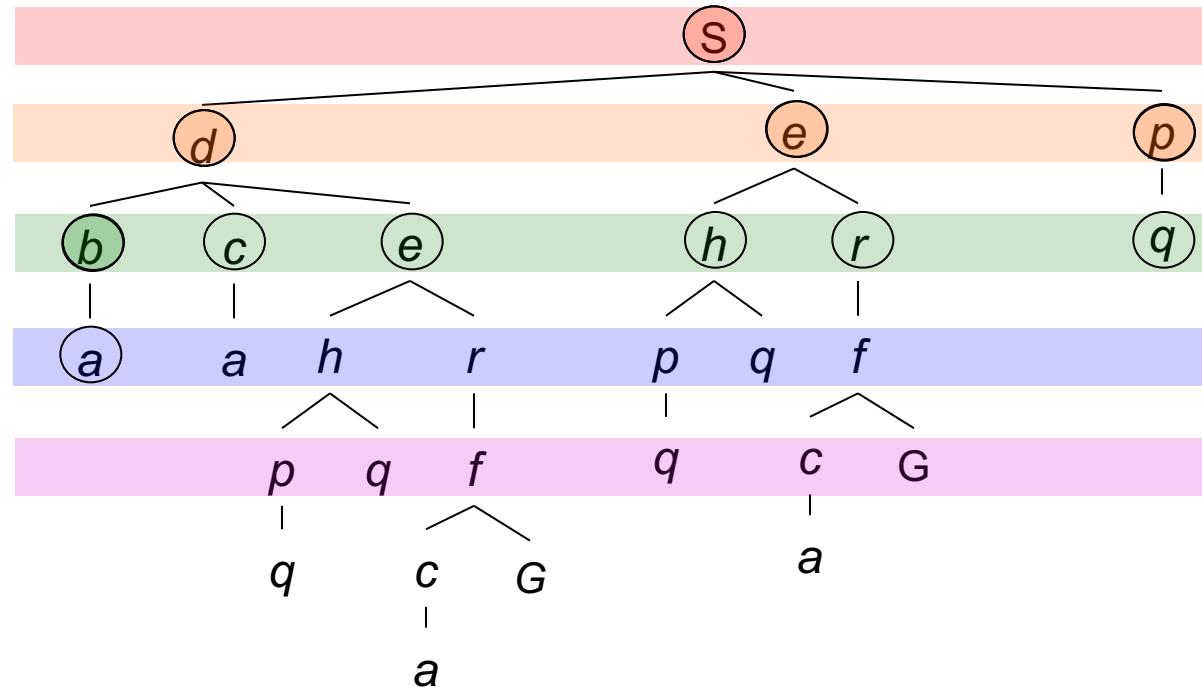
# Breadth-First Search

*Strategy: expand a shallowest node first*

*Implementation: Fringe is a FIFO queue*

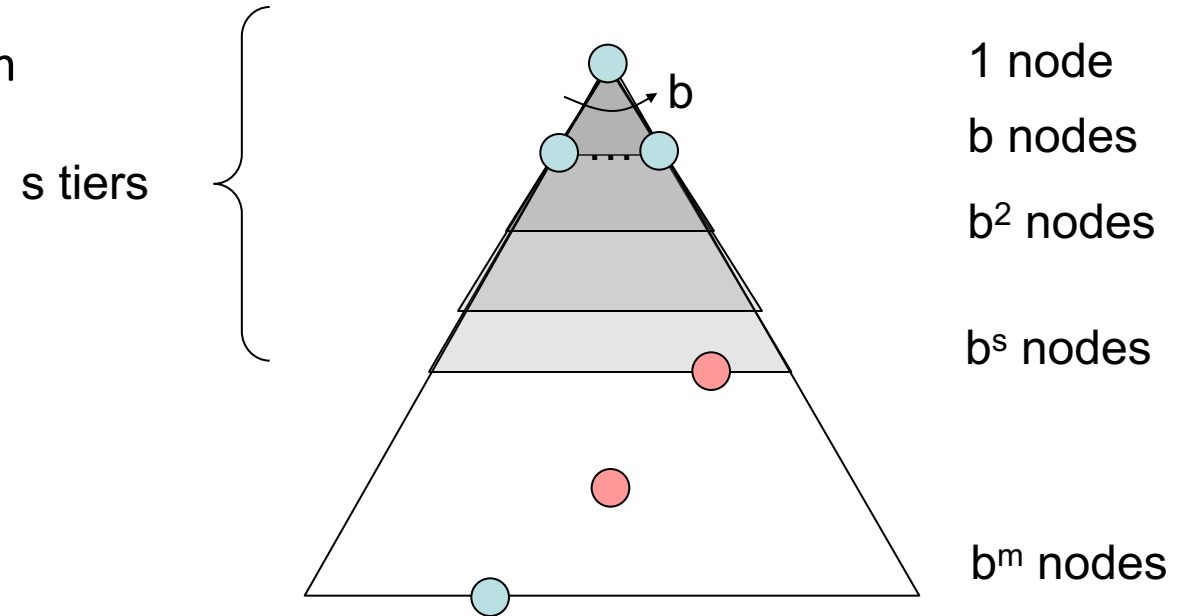


Search  
Tiers

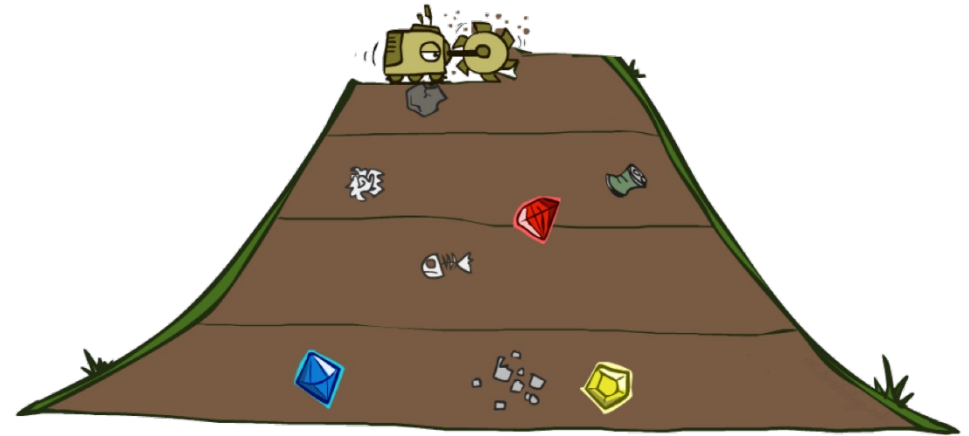


# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Processes all nodes above shallowest solution
  - Let depth of shallowest solution be  $s$
  - Search takes time  $O(b^s)$
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^s)$
- Is it complete?
  - $s$  must be finite if a solution exists, so yes!
- Is it optimal?
  - Only if costs are all 1 (more on costs later)



# Quiz: DFS vs BFS





# Quiz: DFS vs BFS

---

- When will BFS outperform DFS?
- When will DFS outperform BFS?

# Video of Demo Maze Water DFS/BFS (part 1)

---



# Video of Demo Maze Water DFS/BFS (part 2)

---

