

20-08-2024

# Project proposal

- Submit a one page project proposal- deadline 27-08-2024
- Presentation 5 mins - prepare 5 slides
  - ▶ Title slide - project title, group name, name
  - ▶ What?
  - ▶ Why?
  - ▶ How?
  - ▶ Timeline and work division

# Project proposal

- Some pointers:
  - Bengali speech-to-text (RKMVERI data ?)
  - Image based plant disease detection
- Any one interested in web/Apps development ?

# Weights/Parameters updates

- Learning rate based on parameter's value
  - How about use of gradient?
- AdaGrad (Duchi et al., 2011)
  - $A_i = A_i + \left(\frac{\partial L}{\partial w_i}\right)^2; \forall i$
  - $w_i = w_i - \frac{\alpha}{\sqrt{A_i}} \left(\frac{\partial L}{\partial w_i}\right); \forall i$
  - Can you see any problem?
    - Do we need ancient history?
- RMSprop (Hinton class slide, 2012)
  - $A_i = \rho A_i + (1 - \rho) \left(\frac{\partial L}{\partial w_i}\right)^2; \rho \in (0,1); \forall i$
  - $w_i = w_i - \frac{\alpha}{\sqrt{A_i}} \left(\frac{\partial L}{\partial w_i}\right); \forall i$

# Weights/Parameters updates (cont...)

- Adadelta (Zeiler, 2012): replace  $\alpha$  in RMSprop

- ▶  $A_i = A_i + \left( \frac{\partial L}{\partial w_i} \right)^2; \forall i$

- ▶  $\Delta w_i = \frac{\alpha}{\sqrt{A_i}} \left( \frac{\partial L}{\partial w_i} \right); \forall i$

- ▶  $\delta_i = \rho \delta_i + (1 - \rho) \left( \Delta w_i \right)^2; \forall i$

- ▶  $w_i = w_i - \sqrt{\frac{\delta_i}{A_i}} \left( \frac{\partial L}{\partial w_i} \right); \forall i$

- ▶ Can you see any problem?

# Weights/Parameters updates (cont...)

- Adaptive Moment Estimation (Adam- King & Ba, 2025)

- ▶  $m_i = \rho_1 m_i + (1 - \rho_1) \left( \frac{\partial L}{\partial w_i} \right); \rho_1 \in (0,1) ; \forall i$

- ▶  $v_i = \rho_2 v_i + (1 - \rho_2) \left( \frac{\partial L}{\partial w_i} \right)^2 ; \rho_2 \in (0,1) ; \forall i$

- ▶  $\bar{m}_i = \frac{m_i}{1 - \rho_1^t}; \forall i$

- ▶  $\bar{v}_i = \frac{v_i}{1 - \rho_2^t}; \forall i$

- ▶  $w_i = w_i - \alpha \frac{\bar{m}_i}{\sqrt{\bar{v}_i} + \epsilon}; \forall i$

# Weights/Parameters updates (cont...)

- Adaptive Moment Estimation (Adam- Kingma & Ba, 2025)

---

**Algorithm 1:** *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation.  $g_t^2$  indicates the elementwise square  $g_t \odot g_t$ . Good default settings for the tested machine learning problems are  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$  and  $\epsilon = 10^{-8}$ . All operations on vectors are element-wise. With  $\beta_1^t$  and  $\beta_2^t$  we denote  $\beta_1$  and  $\beta_2$  to the power  $t$ .

---

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1)$ : Exponential decay rates for the moment estimates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)

$v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)


**end while**

**return**  $\theta_t$  (Resulting parameters)

---



# Weights/Parameters updates (cont...)







# Convolutional neural network


# LeNet-5

- LeNet-5: 1998
  - #parameters: 61,706



# AlexNet

- Breakthrough in ML, 2012
  - Error: 16.4
  - #parameters: 60,000,000



# VggNet

- VggNet: 2015
  - Error: 7.3%
  - #parameters: 140,000,000

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 <b>LRN</b>	conv3-64 <b>conv3-64</b>	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 <b>conv3-128</b>	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 <b>conv1-256</b>	conv3-256 conv3-256 <b>conv3-256</b>	conv3-256 conv3-256 conv3-256 <b>conv3-256</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Source: Karen Simonyan\* & Andrew Zisserman, ICLR, 2015

# ResNet

- ResNet: 2015
  - Error: 3.6%
  - #parameters: 58,161,162


layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
conv2_x	56×56	3×3 max pool, stride 2				
		$\begin{bmatrix} 3\times 3, 64 \\ 3\times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 64 \\ 3\times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3\times 3, 128 \\ 3\times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 128 \\ 3\times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3\times 3, 256 \\ 3\times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 256 \\ 3\times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3\times 3, 512 \\ 3\times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times 3, 512 \\ 3\times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

ResNet Architectures

Source: Hei et al., Deep Residual Learning for Image Recognition, CVPR 2016


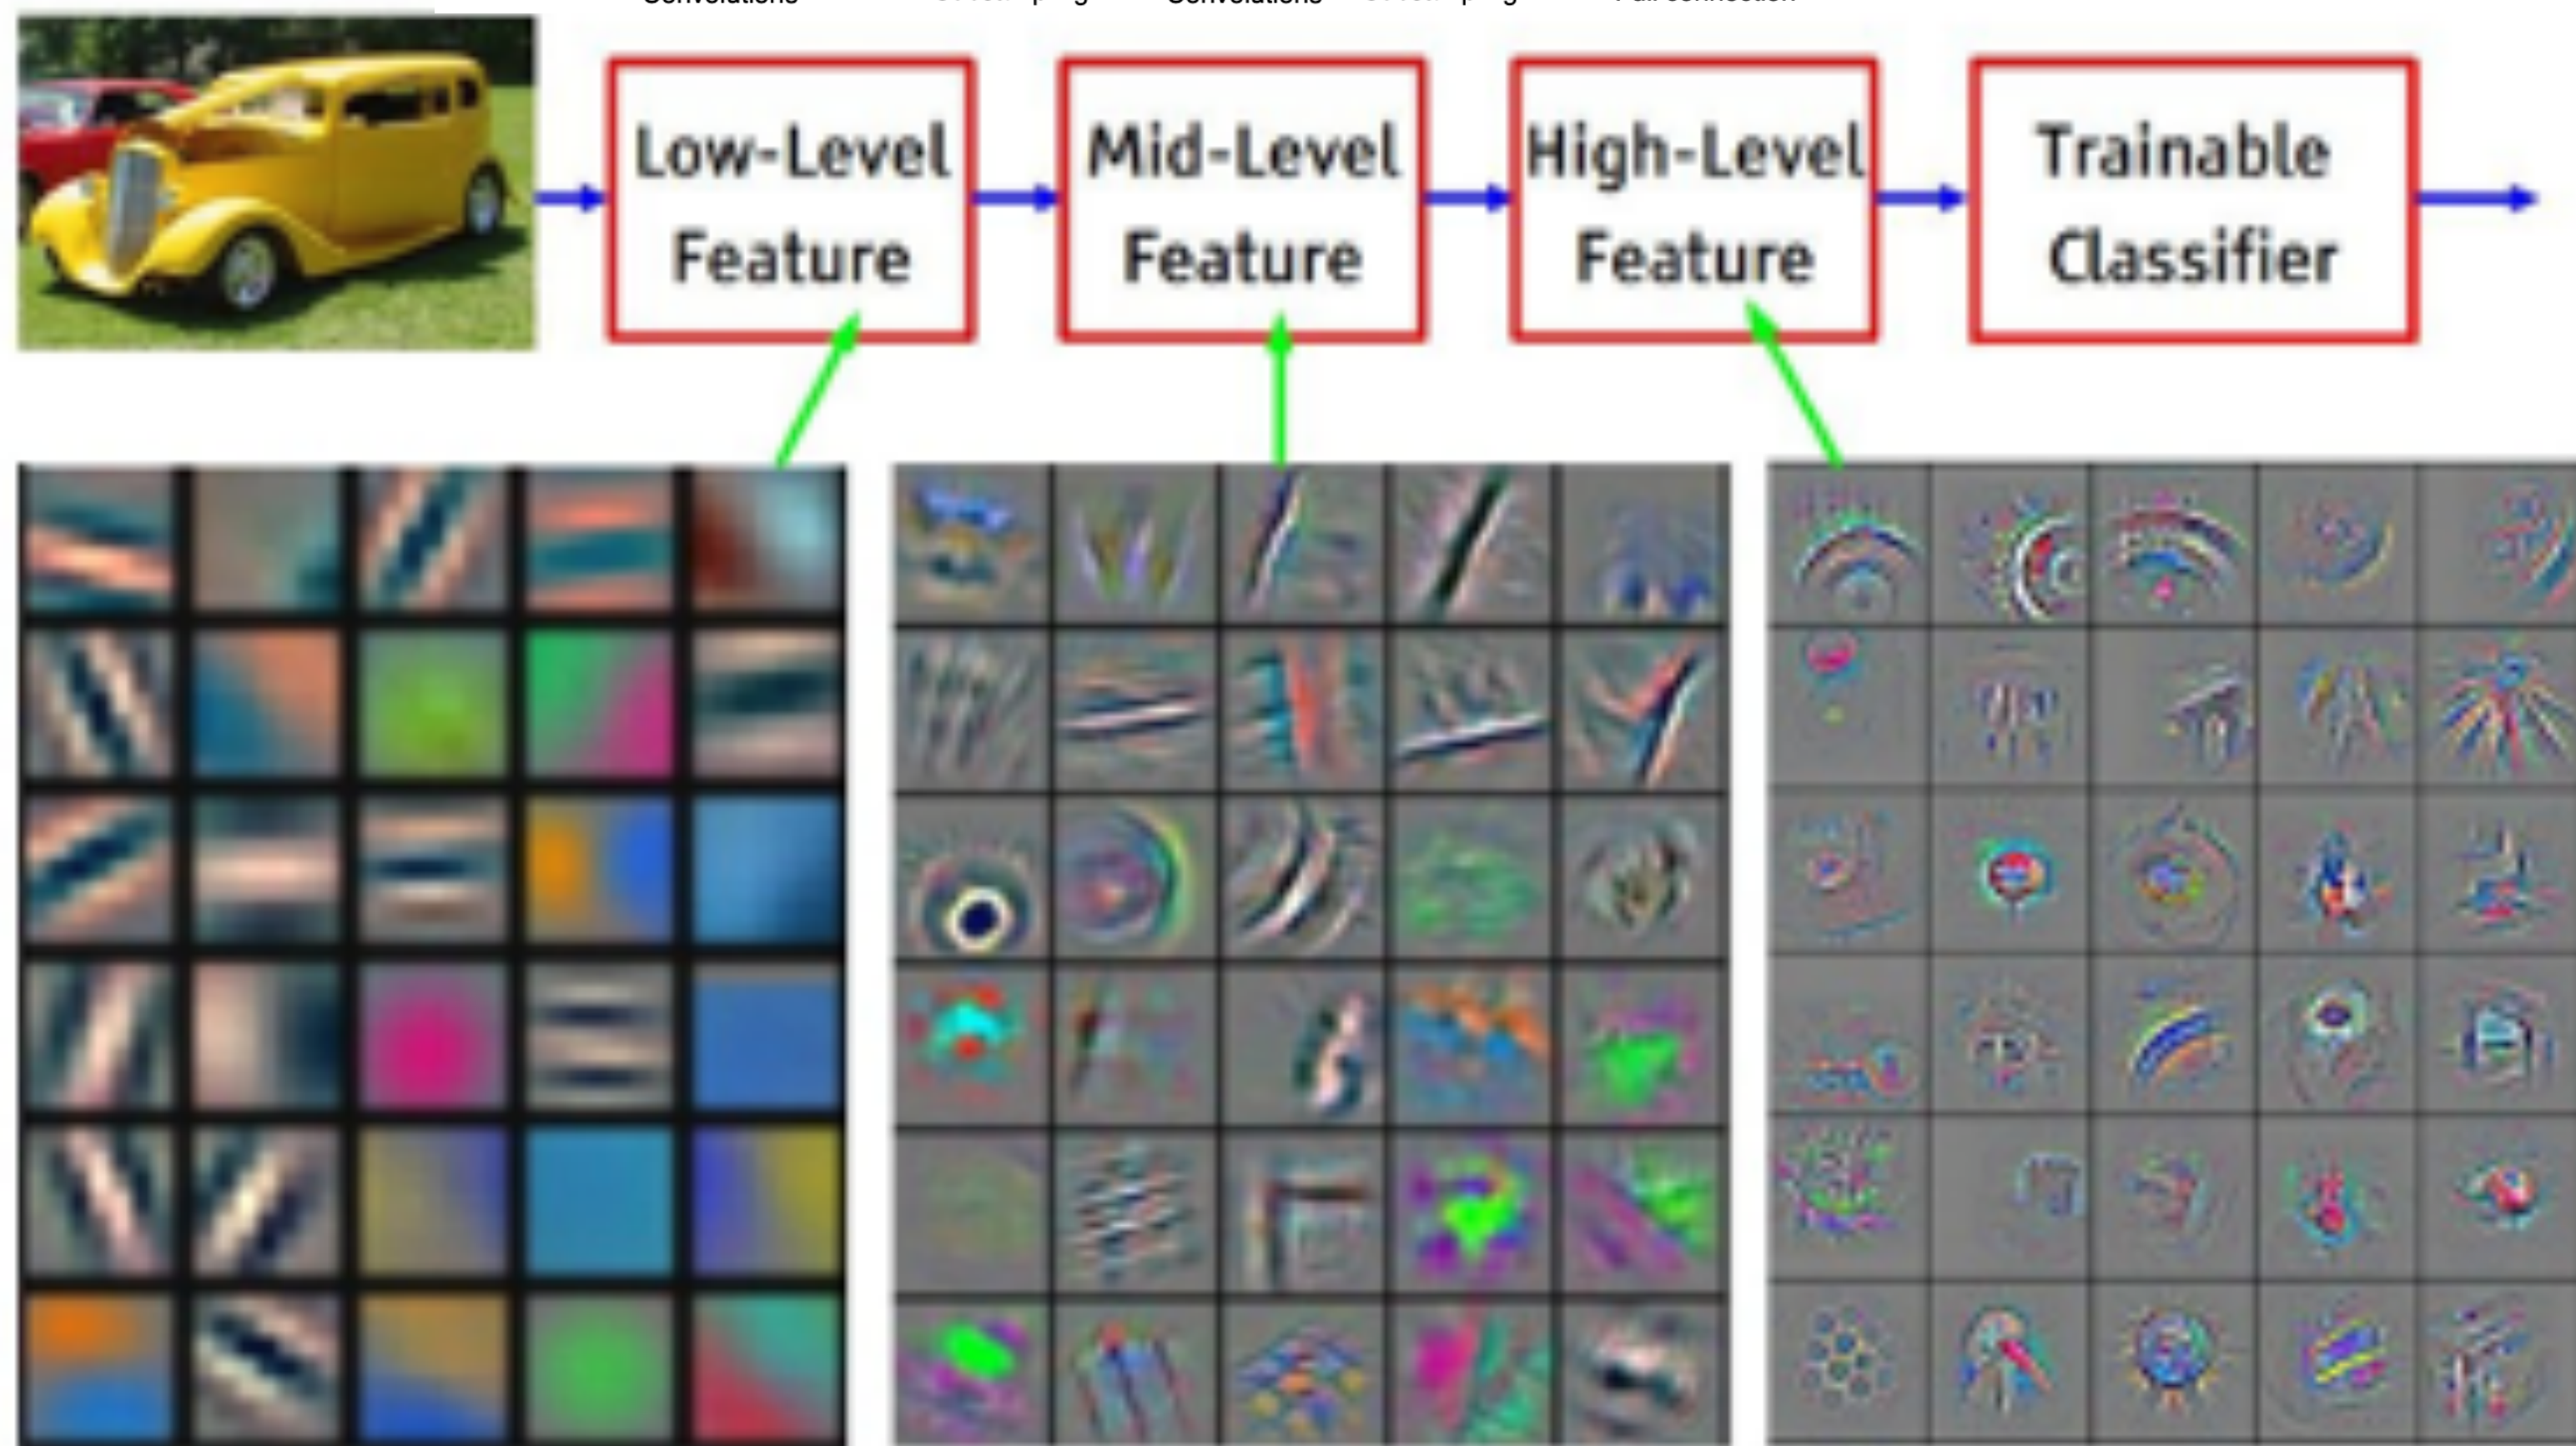
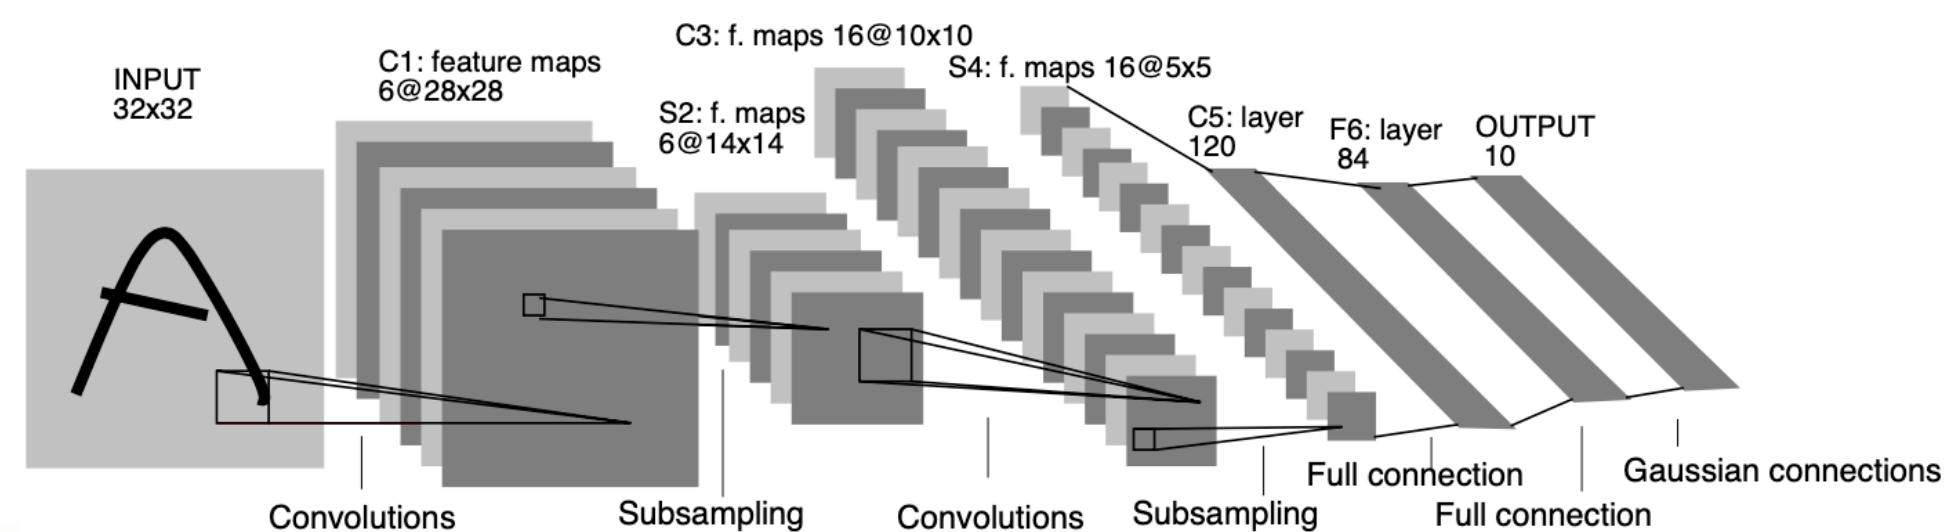


# Results on ImageNet dataset





# Why convolution ?





# Convolution operation

